

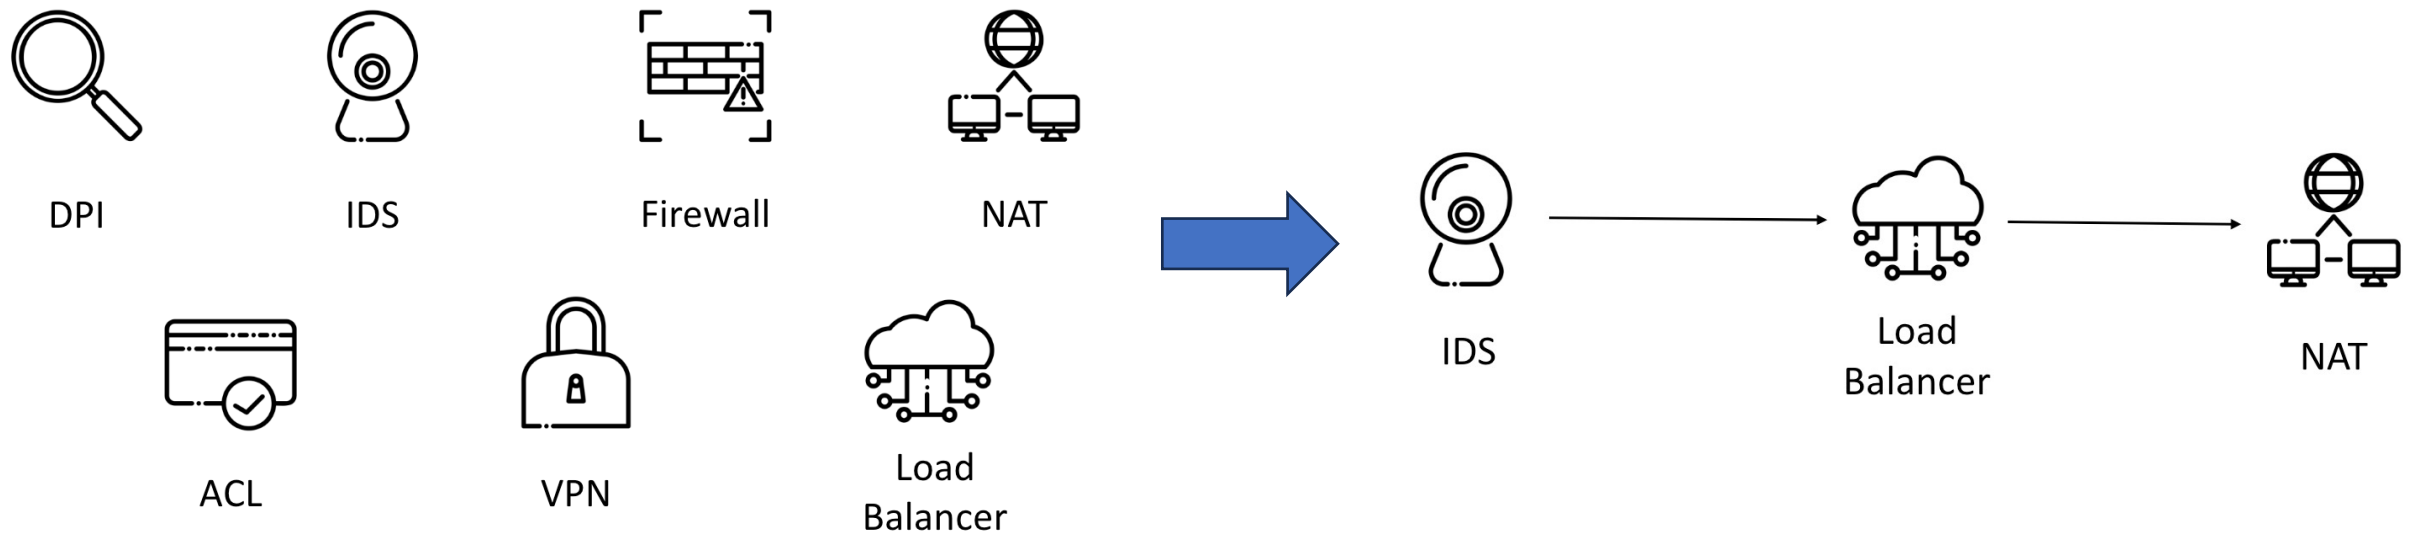
LemonNFV: Consolidating Heterogeneous Network Functions at Line Speed

Hao Li, Yihan Dang et. al. Xian Jiaotong University

NSDI23

VNF and SFC


- Virtualized Network Function (VNF)
- Service Function Chain (SFC)



Difficulties in NFs' Interoperation

- **Heterogeneous NFs Are Not Interoperable**

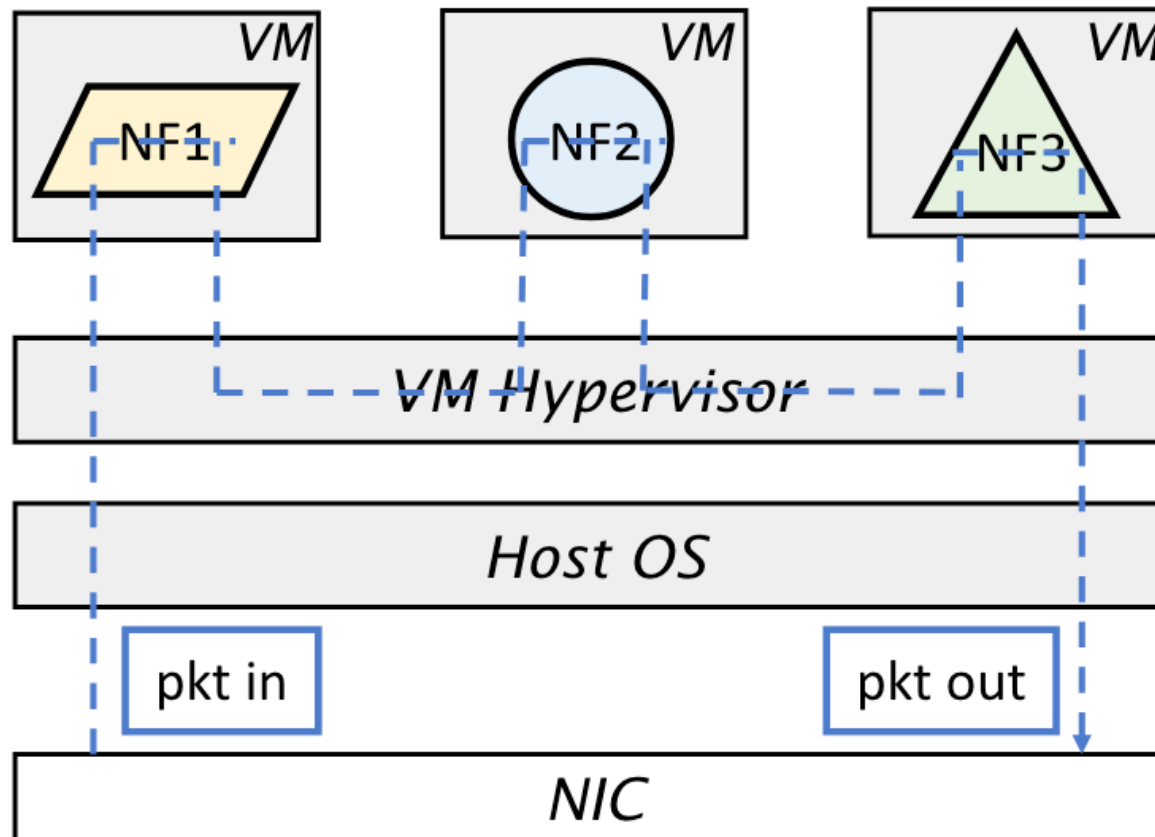
- Programming language
- Execution model
- State & Packet Abstraction



How can heterogeneous NFs interoperate?


Existing Solutions

- Solution 1: Virtualization

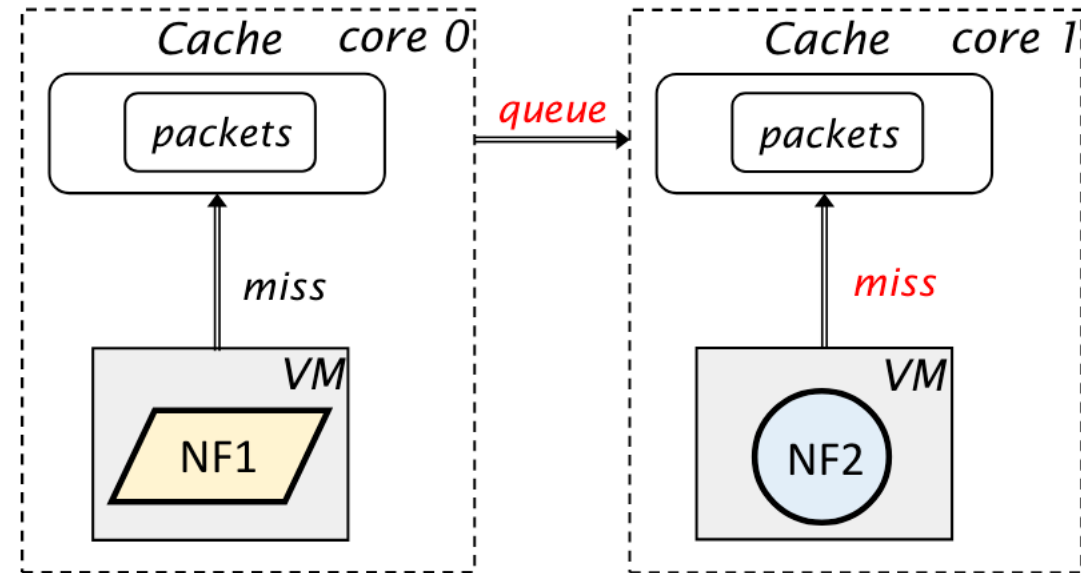


Existing Solutions

- Dilemma: Context Switch vs Inter-Core Traffic



- Scheduling instances on the same core

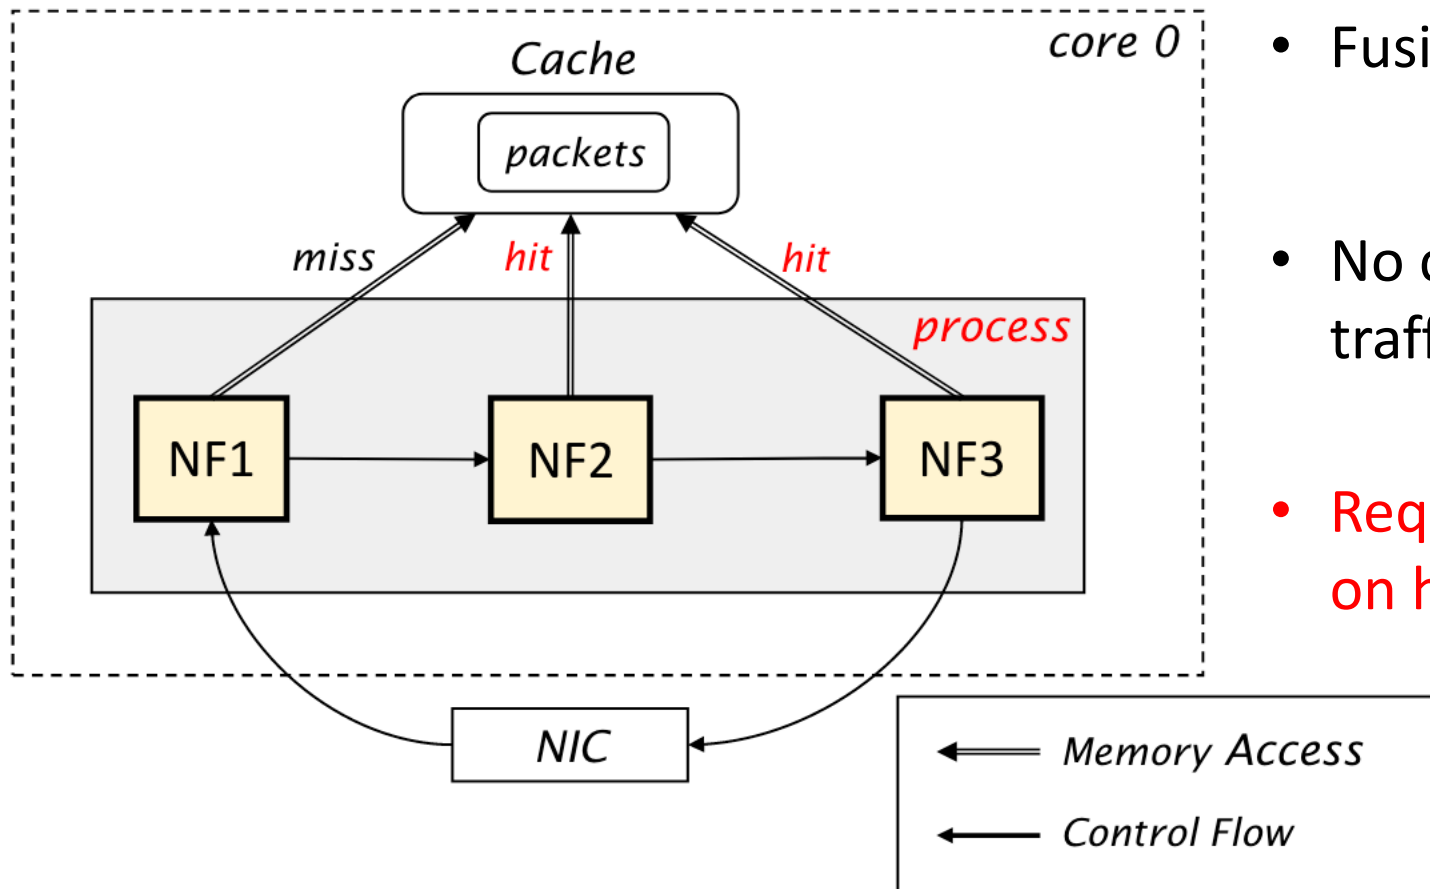


- Pinning instances on dedicated cores

- Virtualization approaches are hard to reach line rate

Existing Solutions

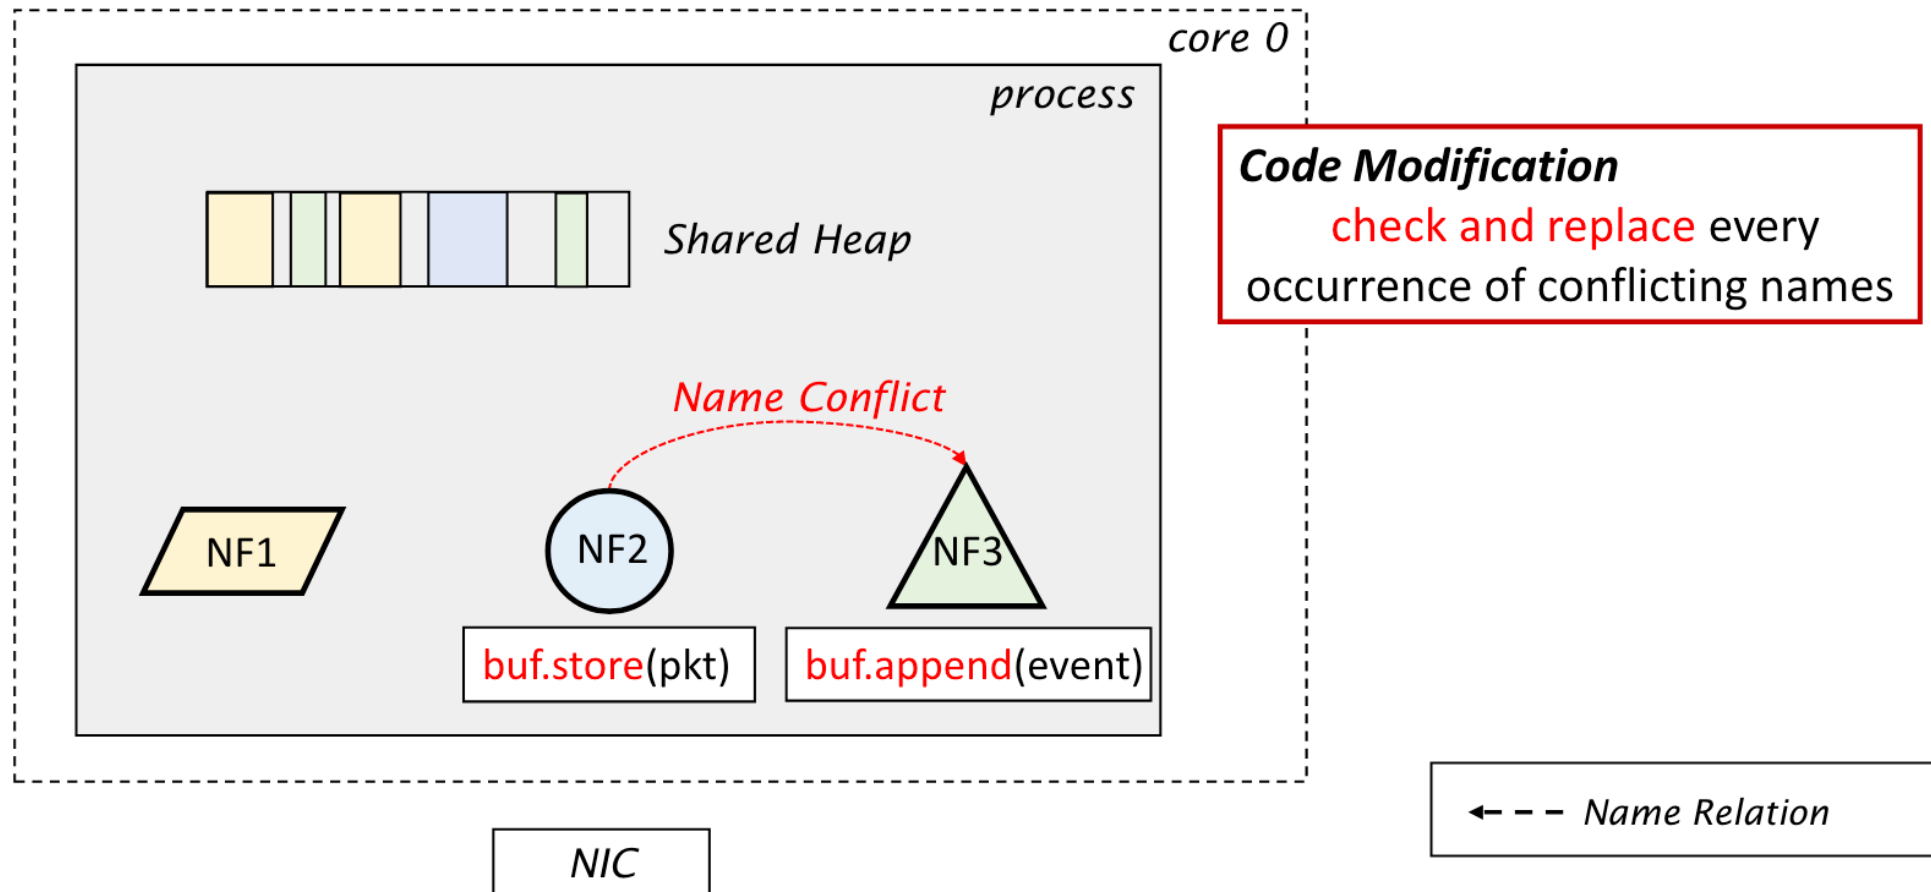
- Solution 2: Consolidation



- Fusing all NFs into one process
- No context switching or inter-core traffic
- Requiring huge code modification on heterogeneous NFs

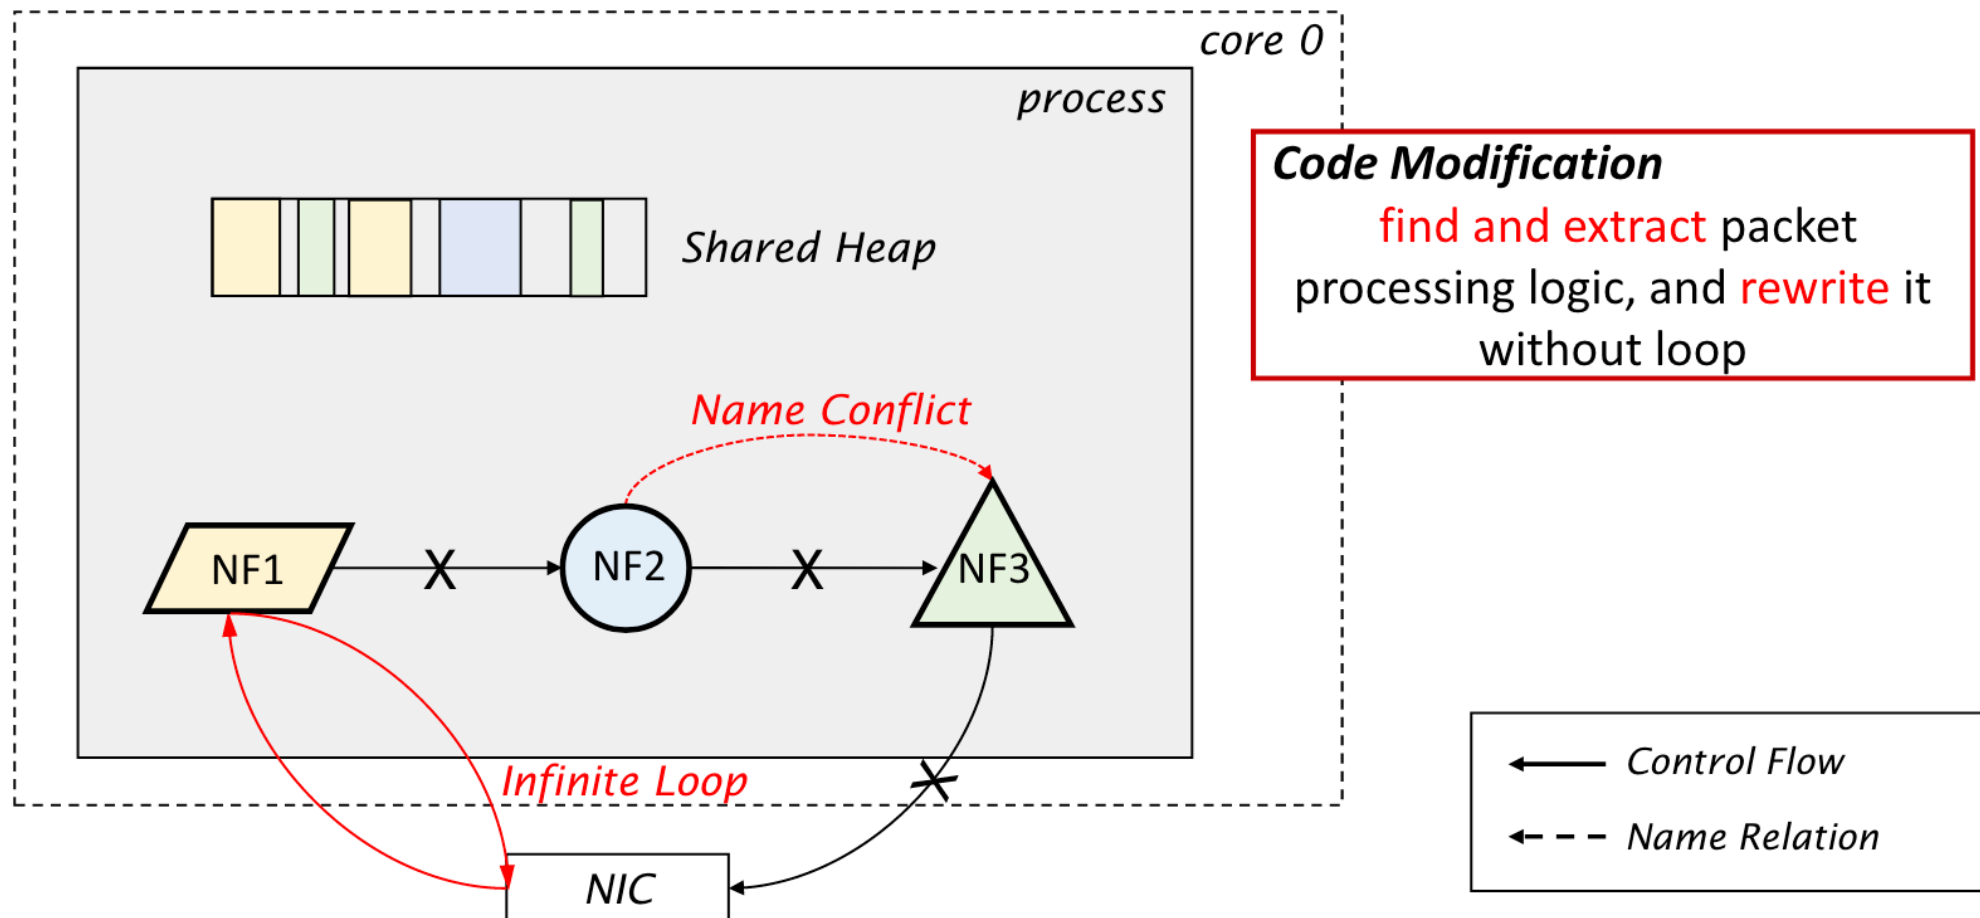
Existing Solutions

- Problem 1: Namespace Conflict



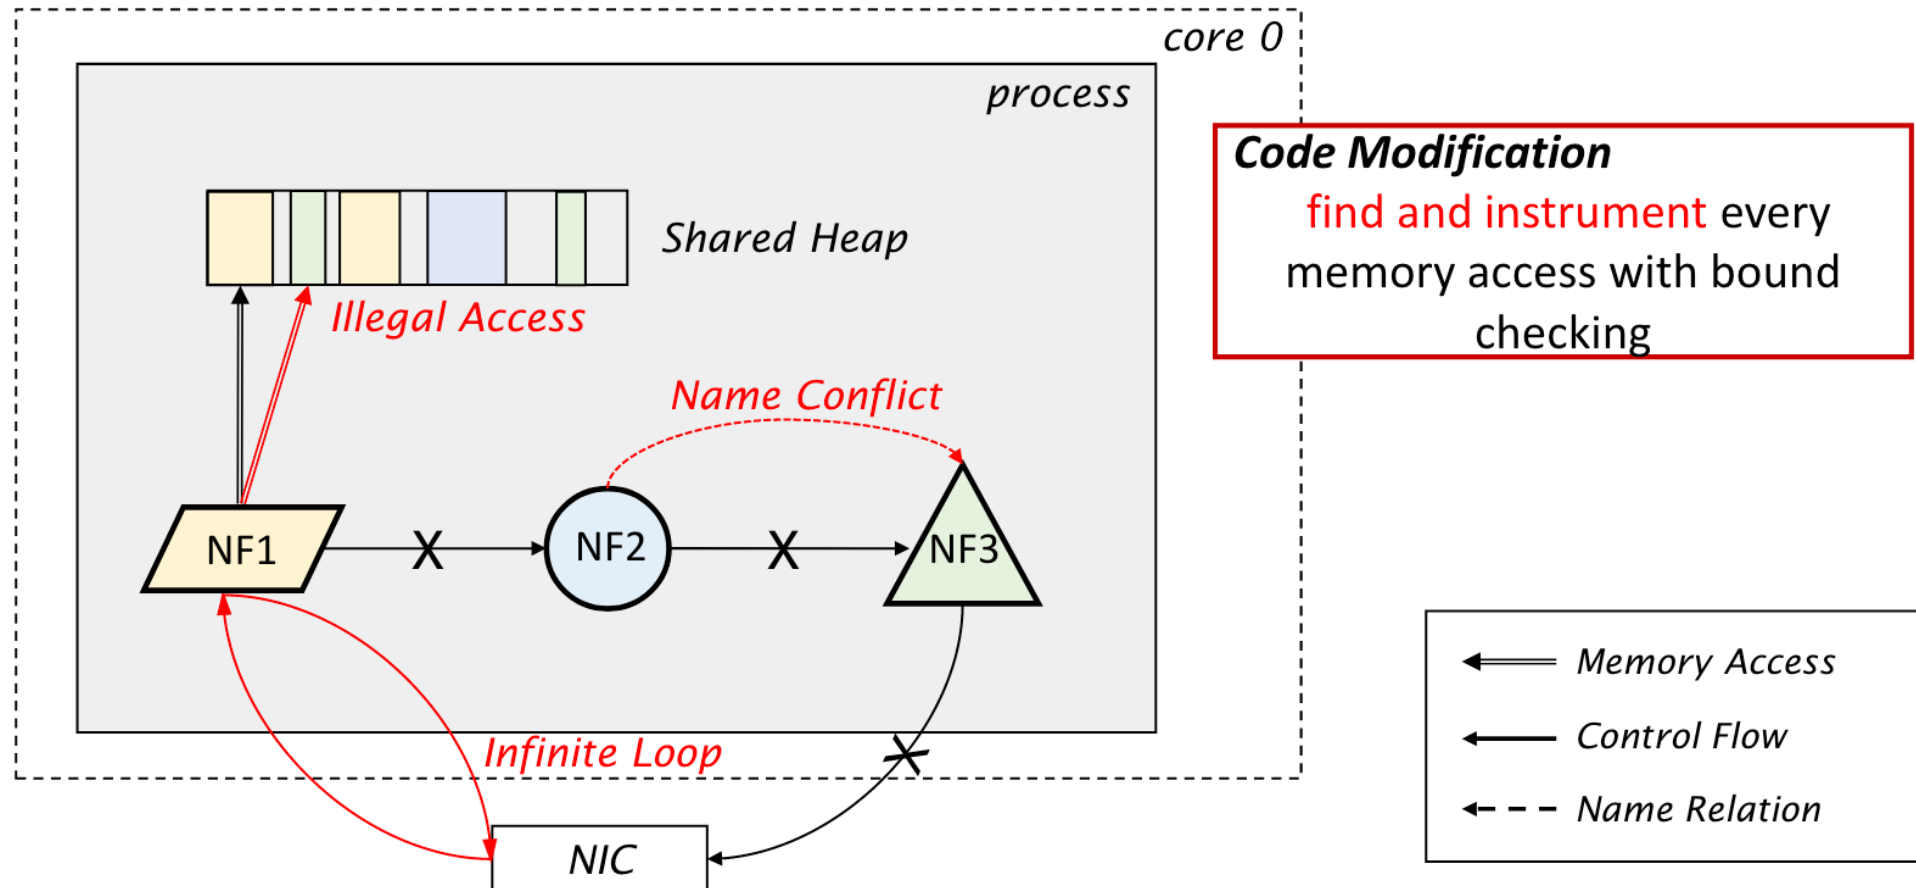
Existing Solutions

- Problem 2: Private Control Flow



Existing Solutions

- Problem 3: Illegal Memory Access

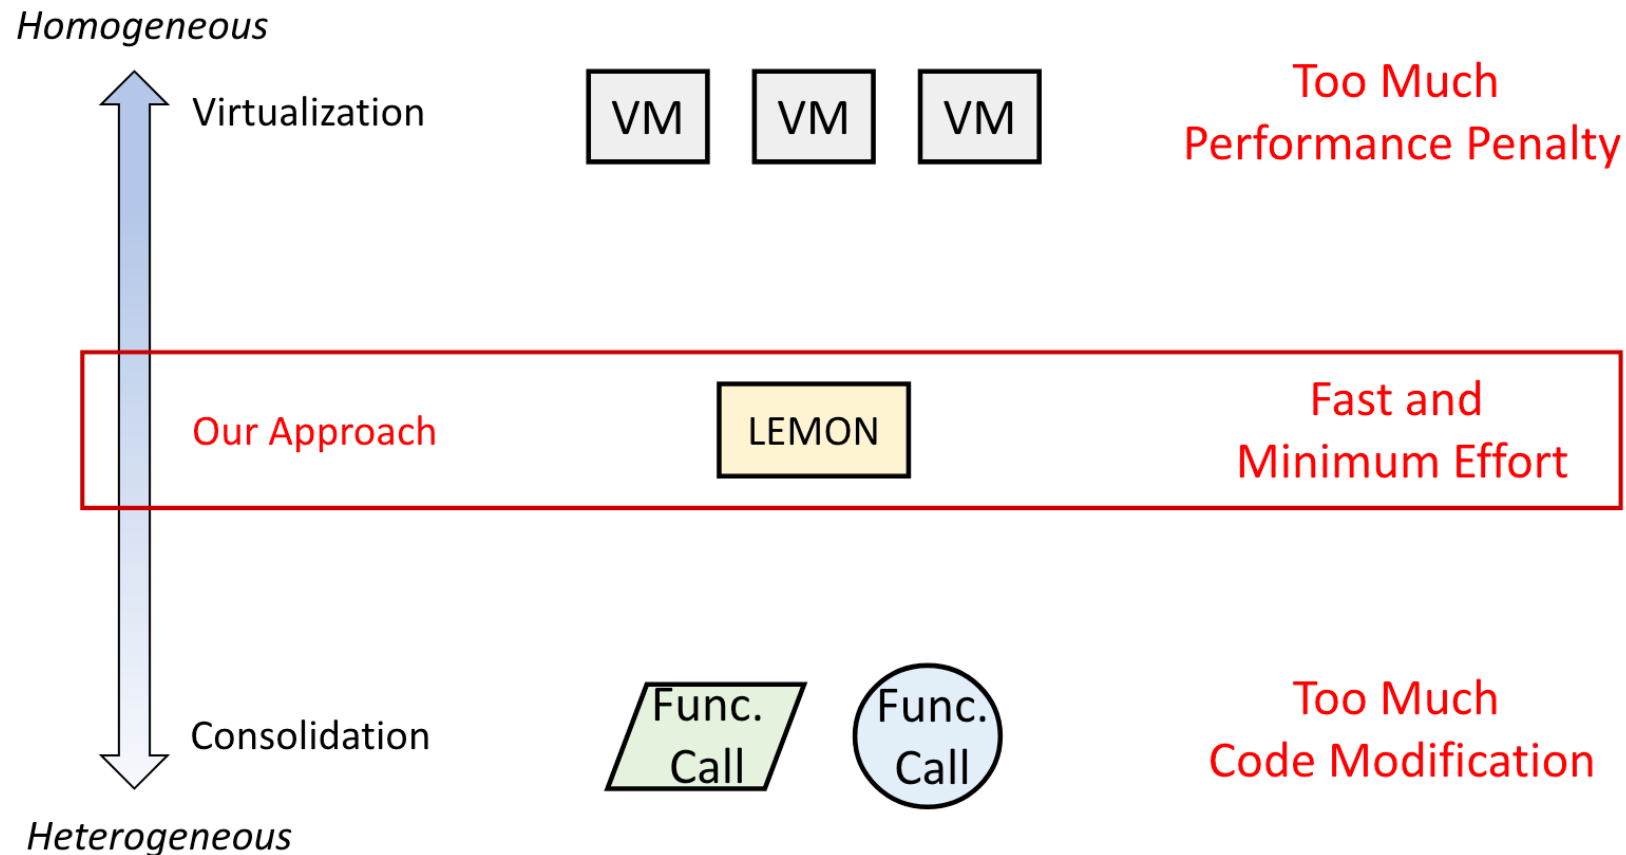


Takeaways on Existing Solutions

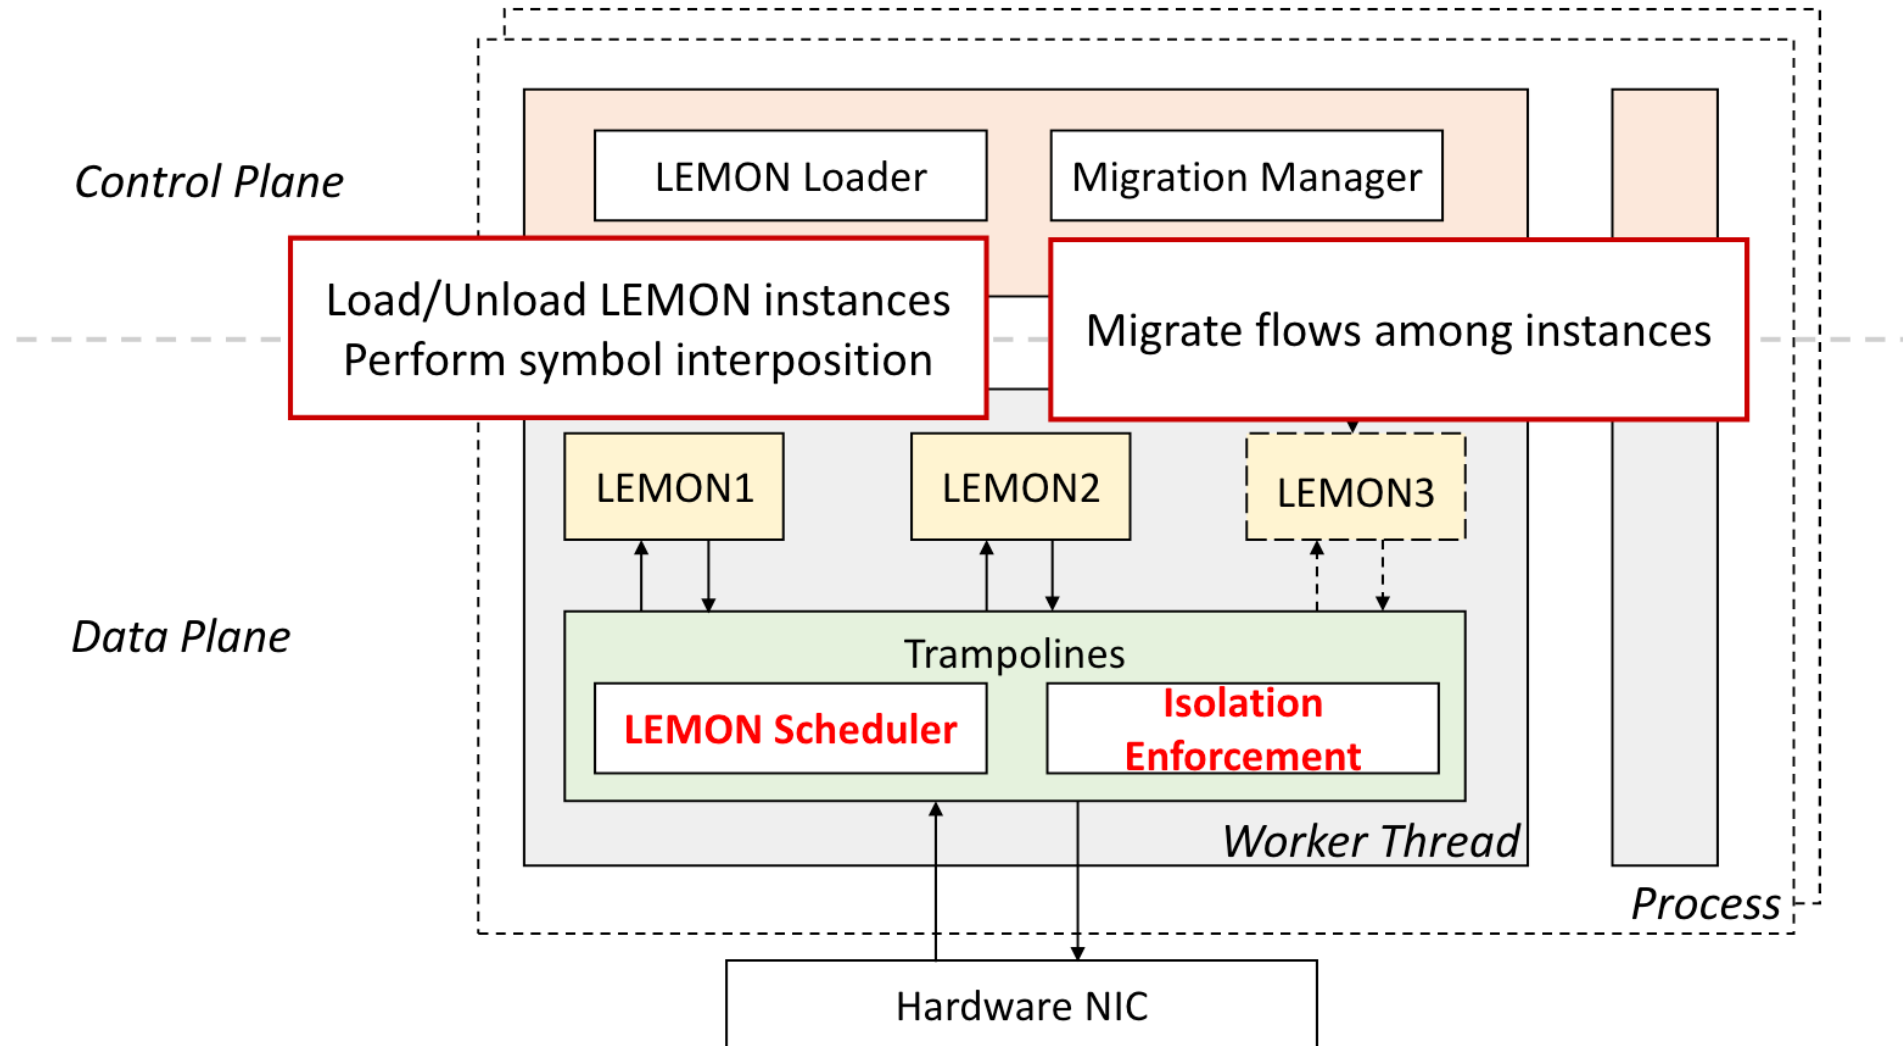
- Solution 1: Virtualization
 - Good: No modification to NFs' codes
 - Bad: Virtualization approaches are hard to reach line rate!
- Solution 2: Consolidation
 - Good: High performance
 - Bad: Direct Consolidation Forces Huge Code Modification!

LemonNFV's Solution

- Insights: exploit the advantages of the two methods

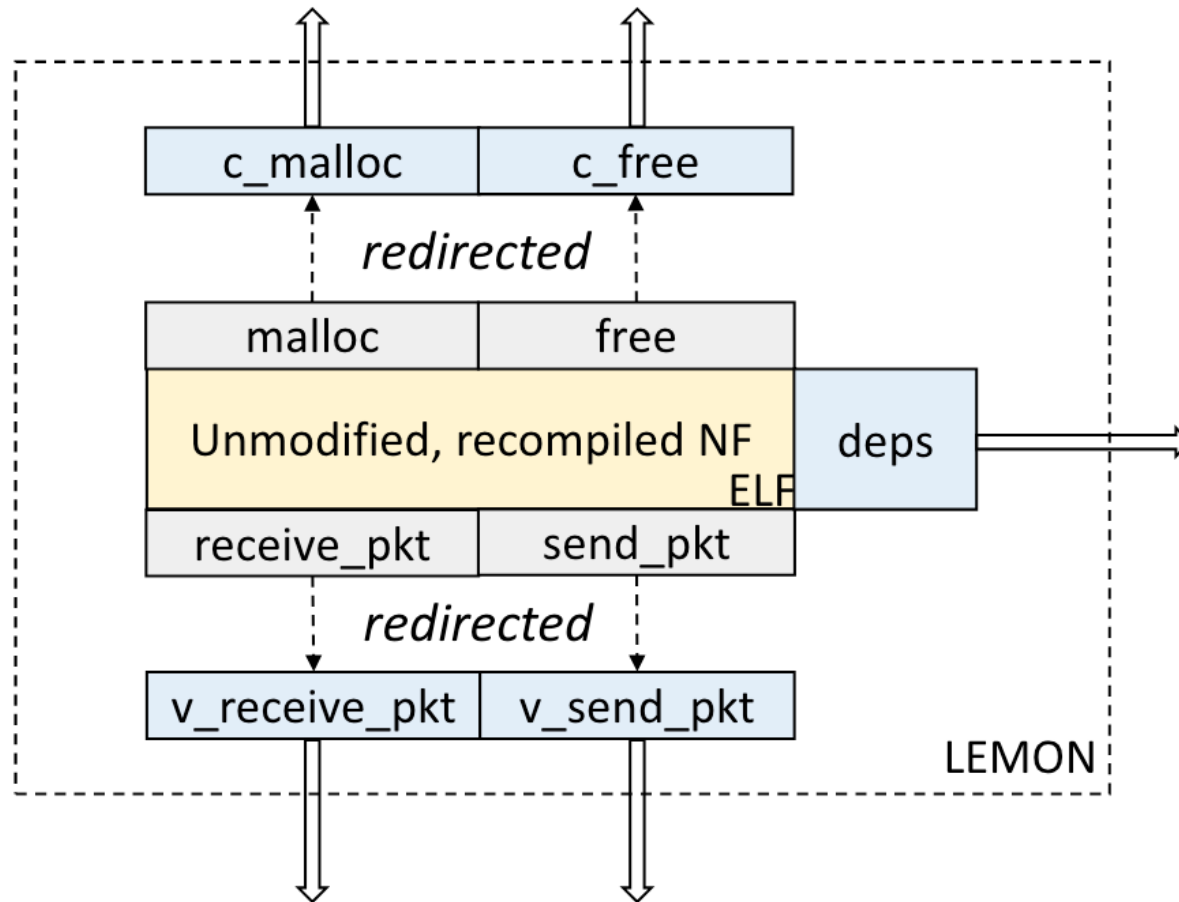


LemonNFV Overview



The LEMON Abstraction (LEast Modified network functiON)

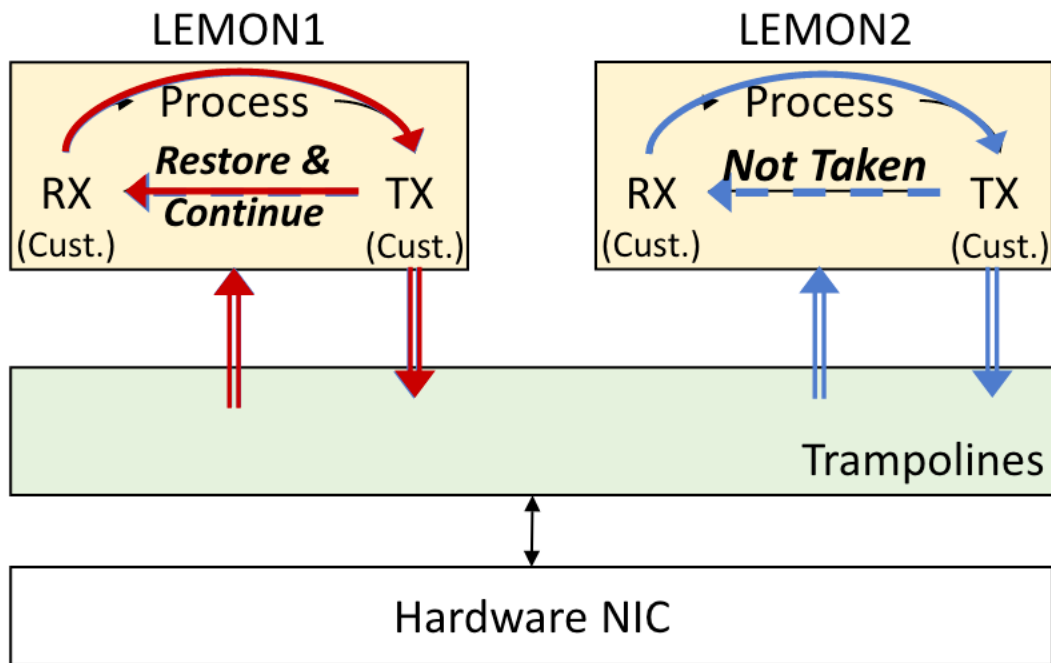
Detect and Prevent Illegal Memory Accesses



Resolve Name Conflicts
Automatically & Correctly

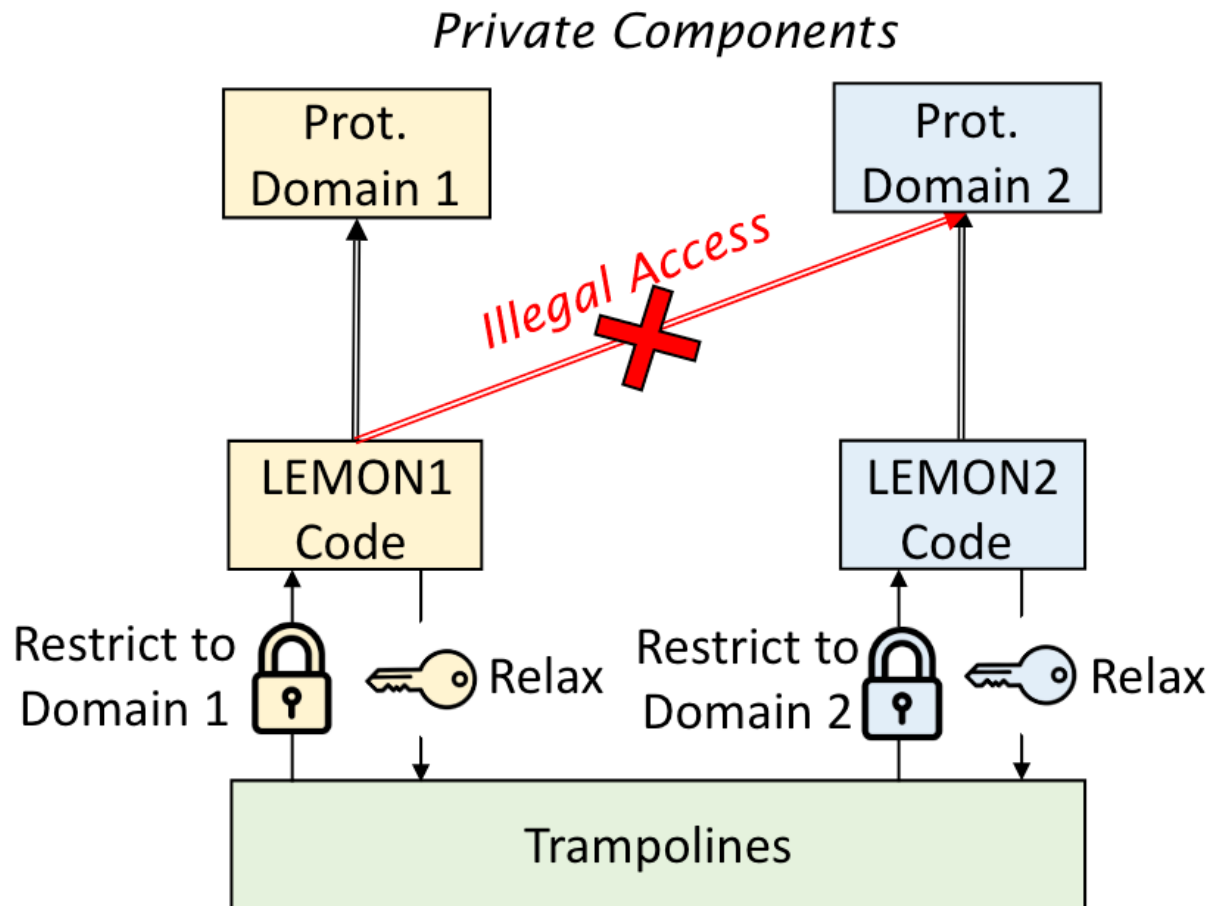
Find and Schedule Infinite Loops

Scheduling the LEMONs with Customized I/O



- By default, NFs process packets in an infinite loop
 - RX/TX talks directly to NIC
- Customized I/O **does not modify the loop** but schedules it
 - Using RX/TX as scheduling points
 - Calling TX **saves context** and returns to the trampolines
 - The trampolines select the next LEMON and **restore its (post-TX) context**


Preventing Illegal Memory Accesses



- The design of LEMON creates bounded memory regions
 - Private heap, stack and dependencies instead of shared ones
 - Accesses outside its own region is illegal
- Bounded memory is efficiently isolated by domain switching
 - LemonNFV uses Intel Protection Key for Userspace (PKU)
 - Restrict access before switching to LEMONS, and relax it before switching back to trampolines

Design Takeaway

Homogeneous





Virtualization



Our Approach



Consolidation




Heterogeneous

Transparent to Users
Memory Isolation



Scheduling and Isolation

Intra-process Execution



High Performance

Evaluation

- Effort of LemonNFV to consolidate heterogeneous NFs
- Performance compared with State-Of-The-Art NFV systems

Evaluation



- Minimum LOC Modification to Interoperation

Heterogeneity of Real World NFs				Huge Code Base Of Real World NFs		Effort of LemonNFV
NF	Framework	Language	I/O	NF LOC	Framework LOC	Modified LOC
IDS	Rubik	C	DPDK	337	31K	2
NAT	FastClick	C++	DPDK	94	331K	2
ACL	NetBricks	Rust	DPDK	401	58K	8
CT	mOS	C	libpcap	325	139K	4
DPI	nDPI	C	libpcap	4498	121K	2

- LemonNFV consolidates heterogeneous NFs **without much effort (LOC)**

Evaluation

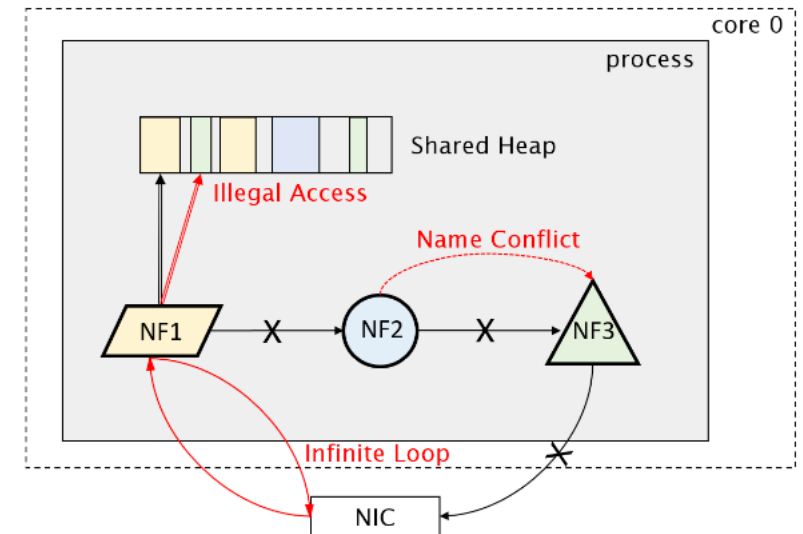


- Comparing Performance with State-Of-The-Art



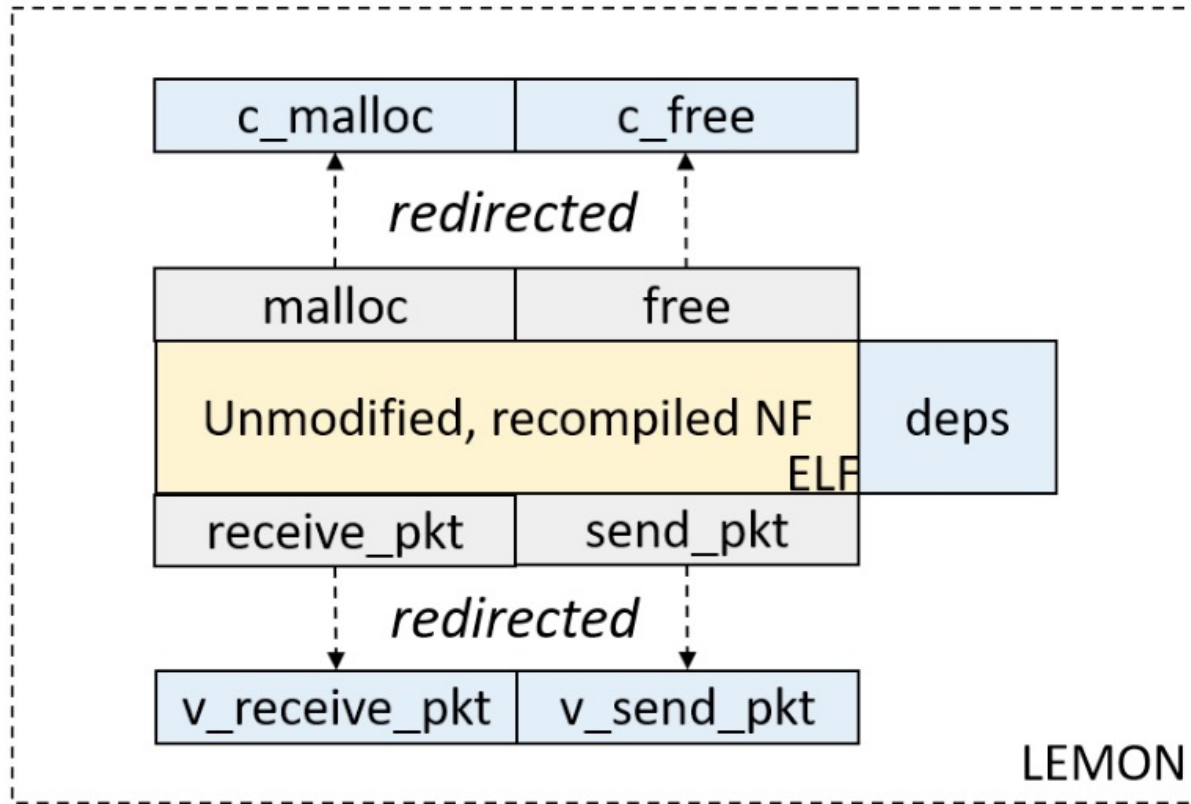
- LemonNFV consolidates heterogeneous NFs with minor overhead

Summary

- Virtualization nor direct consolidation achieves heterogeneous NF interoperation
 - Virtualization overhead
 - Effort of code modification



Summary



- LemonNFV consolidates NFs with minor overhead and effort
 - Designs a unique abstraction LEMON
 - Schedules and isolates LEMONs inside one process