

Shuhao Zhang

Nationality: Singapore

Birth Date: 09/June/1990

August 2023

Information Systems Technology and Design (ISTD) Pillar,
Singapore University of Technology and Design (SUTD),
Singapore.

<https://shuhaozhangtony.github.io/>

+65 6499 7153

shuhao_zhang@sutd.edu.sg

ShuhaoZhangTony

Employment

- 2023– **Assistant Professor**, School of Computer Science and Engineering (SCSE), Nanyang Technological University (NTU).
- 2021–2023 **Assistant Professor**, Information Systems Technology and Design (ISTD), Singapore University of Technology and Design (SUTD).
- 2020–2021 **Postdoc Research Fellow**, Database Systems and Information Management Group, Technische Universität Berlin (TUB).

Consultancy Position

- 2022–2023 **Technical Consultant** (Part-Time), OpenMLDB team, 4paradigm.

Education

- 2019 PhD. Computer Science (Database Systems) National University of Singapore
- 2014 B.Eng. (1st Hons) Computer Engineering Nanyang Technological University, Singapore

Publications

Citations: 714, h-index: 14.

Refereed Research Papers

1. Mao, Y, J Zhao, H Liu, Zhang, Shuhao*, and V Markl (2023). MorphStream: Adaptive Scheduling for Scalable Transactional Stream Processing on Multicores. *Proc. ACM Manag. Data* **1**(1). 26 pages. DOI: 10.1145/3588913.
2. Wang, X, Z Wang, Z Wu, Zhang, Shuhao*, X Shi, and L Lu (2023). Data Stream Clustering: An In-Depth Empirical Study. *Proc. ACM Manag. Data* **1**(2). DOI: 10.1145/3589307.
3. Zeng, X and S Zhang (2023). A Hardware-Conscious Stateful Stream Compression Framework for IoT Applications (Vision). In: *Proceedings of the 17th ACM International Conference on Distributed and Event-Based Systems*. DEBS '23. Neuchatel, Switzerland: Association for Computing Machinery, pp.7–12. DOI: 10.1145/3583678.3596885.
4. Zeng, X and Zhang, Shuhao* (2023). Parallelizing Stream Compression for IoT Applications on Asymmetric Multicores. In: *Proceedings of the 2023 IEEE 39rd International Conference on Data Engineering (ICDE)*.
5. Zhang, H, X Zeng, Zhang, Shuhao, X Liu, M Lu, and Z Zheng (2023). Scalable Online Interval Join on Modern Multicore Processors in OpenMLDB. In: *2023 IEEE 39rd International Conference on Data Engineering (ICDE)*.
6. Zhang, Y, F Zhang, H Li, Zhang, Shuhao, and X Du (2023). CompressStreamDB: Fine-Grained Adaptive Stream Processing without Decompression. In: *Proceedings of the 2023 IEEE 39rd International Conference on Data Engineering (ICDE)*.
7. Meftah, S, Zhang, Shuhao, B Veeravalli, and KMM Aung (2022). Revisiting the Design of Parallel Stream Joins on Trusted Execution Environments. *Algorithms* **15**(6). 20 pages. DOI: 10.3390/a15060183.
8. Xu, Q, F Zhang, M Zhang, J Zhai, B He, C Yang, Zhang, Shuhao, J Lin, H Liu, and X Du (2022). Payment behavior prediction on shared parking lots with TR-GCN. *The VLDB Journal*. 24 pages. DOI: 10.1007/s00778-021-00722-0.
9. Zhang, F, Y Liu, N Feng, C Yang, J Zhai, Zhang, Shuhao, B He, J Lin, X Zhang, and X Du (2021). Periodic Weather-Aware LSTM with Event Mechanism for Parking Behavior Prediction. *IEEE Transactions on Knowledge and Data Engineering*. 14 pages, 1–1. DOI: 10.1109/TKDE.2021.3070202.
10. Zhang, F, C Zhang, L Yang, C Yang, Zhang, Shuhao, B He, W Lu, and X Du (2021). Fine-Grained Multi-Query Stream Processing on Integrated Architectures. *IEEE Transactions on Parallel and Distributed Systems (TPDS)* **32**(9). 14 pages, 2303–2320. DOI: 10.1109/TPDS.2021.3066407.
11. Zhang, Shuhao*, Y Mao, J He, PM Grulich, S Zeuch, B He, RTB Ma, and V Markl (2021). Parallelizing Intra-Window Join on Multicores: An Experimental Study. In: *Proceedings of the 2021 International Conference on Management of Data (SIGMOD)*. SIGMOD '21. 13 pages. Xi'an, Shaanxi, China: Association for Computing Machinery.
12. Zeuch, S, ET Zacharatos, Zhang, Shuhao, X Chatziliadis, A Chaudhary, BD Monte, D Giouroukis, PM Grulich, A Ziehn, and V Markl (2020). NebulaStream: Complex Analytics Beyond the Cloud. *Open Journal of Internet Of Things (OJIOT)* **6**(1). 16 pages, 66–81.

13. Zhang, F, N Feng, Y Liu, C Yang, J Zhai, Zhang, Shuhao, B He, J Lin, and X Du (2020). PewLSTM: Periodic LSTM with Weather-Aware Gating Mechanism for Parking Behavior Prediction. In: *International Joint Conference on Artificial Intelligence(IJCAI)*. 7 pages.
14. Zhang, F, L Yang, Zhang, Shuhao, B He, W Lu, and X Du (2020). FineStream: Fine-Grained Window-Based Stream Processing on CPU-GPU Integrated Architectures. In: *2020 USENIX Annual Technical Conference (USENIX ATC 20)*. 12 pages. USENIX Association, pp.633–647.
15. Zhang, Shuhao, M Plauth, F Eberhardt, A Polze, J Lehmann, G Sejdiu, H Jabeen, L Servadei, C Möstl, F Bär, et al. (2020). *HPI Future SOC Lab-Proceedings 2017*. Universitätsverlag Potsdam.
16. Zhang, Shuhao*, Y Wu, F Zhang, and B He (2020). Towards Concurrent Stateful Stream Processing on Multicore Processors. In: *2020 IEEE 36th International Conference on Data Engineering (ICDE)*. 12 pages, pp.1537–1548. DOI: 10.1109/ICDE48307.2020.00136.
17. Zhang, Shuhao*, F Zhang, Y Wu, B He, and P Johns (2020). Hardware-Conscious Stream Processing: A Survey. *SIGMOD Rec.* **48**(4). 12 pages, 18–29. DOI: 10.1145/3385658.3385662.
18. Ang, J, T Fu, J Paul, Zhang, Shuhao*, B He, TSD Wenceslao, and SY Tan (2019). TraV: An Interactive Exploration System for Massive Trajectory Data. In: *2019 IEEE Fifth International Conference on Multimedia Big Data (BigMM)*. 4 pages, pp.309–313. DOI: 10.1109/BigMM.2019.000–4.
19. Zhang, Shuhao*, J He, AC Zhou, and B He (2019). BriskStream: Scaling Data Stream Processing on Shared-Memory Multicore Architectures. In: *Proceedings of the 2019 International Conference on Management of Data (SIGMOD)*. SIGMOD '19. 18 pages. Amsterdam, Netherlands: Association for Computing Machinery, pp.705–722. DOI: 10.1145/3299869.3300067.
20. Zhang, F, J Zhai, B He, Zhang, Shuhao, and W Chen (2017). Understanding Co-Running Behaviors on Integrated CPU/GPU Architectures. *IEEE Transactions on Parallel and Distributed Systems* **28**(3). 14 pages, 905–918. DOI: 10.1109/TPDS.2016.2586074.
21. Zhang, Shuhao*, B He, D Dahlmeier, AC Zhou, and T Heinze (2017). Revisiting the Design of Data Stream Processing Systems on Multi-Core Processors. In: *2017 IEEE 33rd International Conference on Data Engineering (ICDE)*. 12 pages, pp.659–670. DOI: 10.1109/ICDE.2017.119.
22. Zhang, Shuhao*, HT Vo, D Dahlmeier, and B He (2017). Multi-Query Optimization for Complex Event Processing in SAP ESP. In: *2017 IEEE 33rd International Conference on Data Engineering (ICDE)*. 12 pages, pp.1213–1224. DOI: 10.1109/ICDE.2017.166.
23. Tang, S, B He, Zhang, Shuhao, and Z Niu (2016). Elastic Multi-resource Fairness: Balancing Fairness and Efficiency in Coupled CPU-GPU Architectures. In: *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*. 12 pages, pp.875–886. DOI: 10.1109/SC.2016.74.
24. Wang, Z, Zhang, Shuhao, B He, and W Zhang (2016). Melia: A MapReduce Framework on OpenCL-Based FPGAs. *IEEE Transactions on Parallel and Distributed Systems* **27**(12). 14 pages, 3547–3560. DOI: 10.1109/TPDS.2016.2537805.
25. Zhang, F, J Zhai, W Chen, B He, and Zhang, Shuhao (2015). To Co-run, or Not to Co-run: A Performance Study on Integrated Architectures. In: *2015 IEEE 23rd International Symposium on Modeling, Analysis, and Simulation of Computer and Telecommunication Systems*. 4 pages, pp.89–92. DOI: 10.1109/MASCOTS.2015.27.
26. He, J, Zhang, Shuhao, and B He (2014). In-Cache Query Co-Processing on Coupled CPU-GPU Architectures. *Proceedings of the VLDB Endowment* **8**(4). 12 pages, 329–340. DOI: 10.14778/2735496.2735497.
27. Zhang, Shuhao, J He, B He, and M Lu (2013). OmniDB: Towards Portable and Efficient Query Processing on Parallel CPU/GPU Architectures. *Proceedings of the VLDB Endowment* **6**(12). 4 pages, 1374–1377. DOI: 10.14778/2536274.2536319.

Awards/Patents

- 2018-2019 Research Achievement Award, School of Computing, National University of Singapore
- 2017-2018 Research Achievement Award, School of Computing, National University of Singapore
- 2019/12/31 Efficient execution of data stream processing systems on multi-core processors, US Patent (10521432)
- 2018/4/24 Multi-query optimizer for complex event processing, US Patent (9953056)

Service

I have helped in organizing various events.

Years	Event Name	Role
2023	National Research Foundation - Foundational Research Capabilities Study (<i>NRF-FRC, Singapore</i>)	Area Leader
2023	National Supercomputing Centre HPC AI Innovation Challenge 2023 (<i>NSCC, Singapore</i>)	Organising Committee
2022	Singapore Data Science Consortium (<i>SDSC, Singapore</i>) PhD Dissertation Fellowship Evaluation	Evaluator

I have regularly served on the program committee of international conferences such as SC, ICDCS, ICPP, and KDD.

Years	Conference Name	Role
2024	IEEE International Conference on Data Engineering (<i>ICDE</i>)	Chair of TKDE Posters Track & PC Member
2023	NRF FRC Green Computing - Green Edge Workshop	Co-chair
2021-2023	International Conference for High Performance Computing, Networking, Storage, and Analysis (<i>SC</i>)	Program Committee Member
2021-2023	International Conference on Parallel Processing (<i>ICPP</i>)	Program Committee Member
2022-2023	ACM SIGKDD Conference on Knowledge Discovery and Data Mining (<i>KDD</i>)	Program Committee Member
2023	International Conference on Extending Database Technology (<i>EDBT</i>)	Program Committee Member
2023	International Conference on Advanced Data Mining and Applications (<i>ADMA</i>)	Program Committee Member
2020-2021, 2023	IEEE International Conference on High Performance Computing, Data, and Analytics (<i>HiPC</i>)	Program Committee Member
2022	International Conference on Distributed and Event-based Systems (<i>DEBS</i>)	Program Committee Member
2022	Database Systems for Advanced Applications (<i>DASFAA</i>)	Program Committee Member
2020-2021	International Conference on Distributed Computing Systems (<i>ICDCS</i>)	Program Committee Member
2021	ACM Asia-Pacific Workshop on Systems (<i>Apsys</i>)	Program Committee Member
2021	International Symposium on Cluster, Cloud and Internet Computing (<i>CCGrid</i>)	Program Committee Member
2019	ACM International Conference on Information and Knowledge Management (<i>CIKM</i>)	Program Committee Member
2019	IEEE International Conference on Cloud Computing and Intelligence Systems (<i>CCIS</i>)	Program Committee Member

I regularly served as a reviewer for various peer-reviewed journals:

- IEEE Transactions on Knowledge and Data Engineering (TKDE)
- IEEE Transactions on Cloud Computing (TCC)
- ACM Transactions on Database Systems (TODS)
- ACM Transactions on Parallel Computing
- IEEE Transactions on Parallel and Distributed Systems (TPDS)
- IEEE Transactions on Services Computing (TSC)
- Information Systems (IS)
- Pervasive and Mobile Computing (PMC)
- Journal of Computer Science and Technology (JCST)
- Future Generation Computer Systems (FGCS)
- Computers & Security (CS)

I have also served as a proposal evaluator for research grants:

- Future Communications Programme, Singapore (Proposal of 3 Years)
- Israel Science Foundation Personal Research Grants, Israel (Proposal of 4 Years)

Referees

The following referees will email the recommendation letters directly if asked:

- **Prof. Bingsheng He** (PhD Supervisor) - National University of Singapore
Email: hebs@comp.nus.edu.sg
- **Prof. Volker Markl** (Postdoc Advisor) - Technische Universität Berlin
Email: volker.markl@tu-berlin.de
- **Prof. Xiaoyong Du** (Collaborator) - Renmin University of China
Email: duyong@ruc.edu.cn

Appendix

Research income

I have acquired (in most cases solely) about S\$1.6 million in external research grants since 2021. I am a Principal Investigator for a current Singapore MOE Tier 2 Grant and a Principal Investigator for two current Singapore National Research Foundation Grants. I also have research funding from Temasek Laboratories. Details are listed below. * and # indicate the PI and Co-PI, respectively.

2023-2026	Shuhao Zhang* , Mian Lu, Tony Quek. "IntelliStream: Towards Highly-Optimized, Ultra-Scalable, Self-adaptive Data Streaming Analytics in the Heterogeneous Multicore IoT Systems". <i>Funding from Singapore Ministry of Education (MOE) Academic Research Fund (AcRF) Tier 2</i>	S\$500k
2022-2025	Shuhao Zhang* , Binbin Chen#. "A Stream Processing based NFV Platform for 5G on Modern Multicore Processors". <i>Funding from National Research Foundation, Singapore and Infocomm Media Development Authority under its Future Communications Research & Development Programme</i>	S\$496k
2022-2025	Shuhao Zhang* , Mian Lu. "Energy-efficient, Scalable, and Reliable Distributed Green Streaming Machine Learning for Edges". <i>Funding from National Research Foundation, Singapore and Infocomm Media Development Authority under its Future Communications Research & Development Programme</i>	S\$496k
2022-2025	Meixia Lin*, Das Bikramjit#, Wei Quin Yow#, Shuhao Zhang# . "Towards Co-clustering in Big Data: An Optimization Perspective". <i>Funding from SUTD Kickstarter Initiative (SKI)</i>	S\$477k
2023	Shuhao Zhang* , Chun Wei Seah, Wei Lu. "Towards Online Continual Pre-Trained Language Model Maintenance". <i>Funding from TL@SUTD</i>	S\$100k
2022	Shuhao Zhang* , Wei Lu. "Online Sentiment Learning of Massive Data Streams". <i>Funding from TL@SUTD</i>	S\$67k
2022-2025	Shuhao Zhang* . "Revisiting the Algorithms for Clustering Evolving Trajectory Streams". <i>Funding from SUTD-ZJU (VP)</i>	S\$80k
2021-2024	Shuhao Zhang* . "Efficient Intra-Window Join on the Multicore IoT systems". <i>Funding from START-UP RESEARCH GRANT (SRG)</i>	S\$100k

Teaching

Courses Taught

Course	Period	Typical class sizes
50.049 Parallel Computing on Multicore Architectures (Sole Teaching) ¹	Spring 2022	20
50.003 Elements of Software Construction (Co-Teaching)	Spring 2021, Spring 2022, Summer 2023	200+

Recent Student Teaching Feedback

Sem.	Course	Class Size	#Responses	Subject Rating (/5)	Insturctor Rating (/5)
22S1	50.049	17	17	4.5	4.5
23S2	50.003	212	199	3.8	4.1
22S1	50.003	207	192	3.4	3.9
21S1	50.003	172	167	2.7 ²	3.2

¹A course developed by me.

²Mostly due to course management issues. I was assigned to co-teach this module literally two days after I joined the university.