

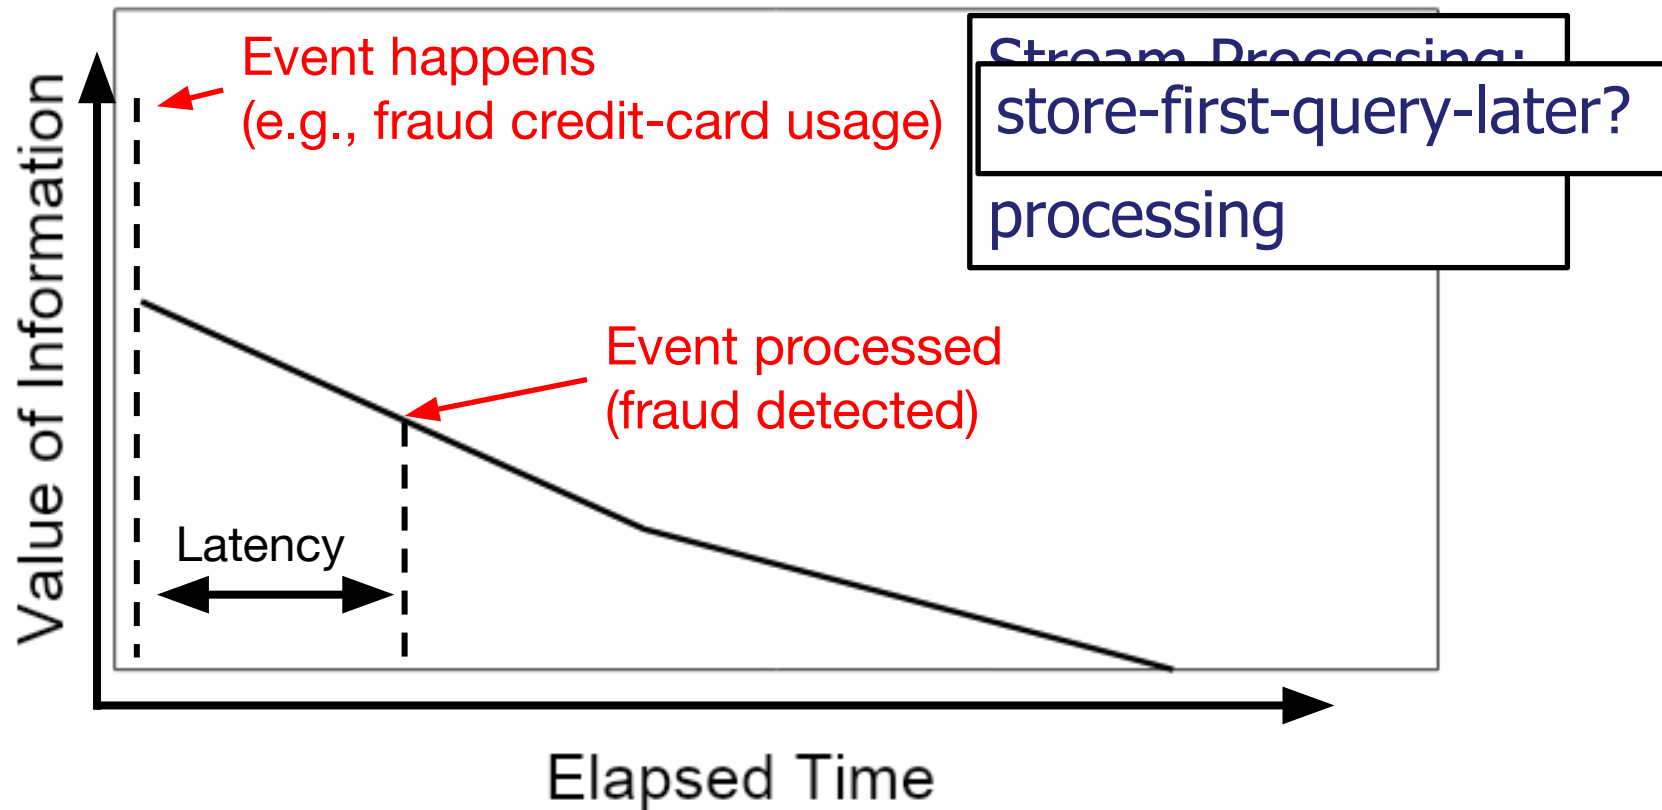
High Performance Stream Processing on Modern Hardware: The Past and The Future

Shuhao Zhang
Postdoc Researcher, DIMA TU Berlin

Outline

- Motivation
 - Why stream processing?
 - Why stream processing on modern HW?
- Past Work
- Future Work

流计算(stream processing)




流计算在阿里巴巴的应用

source: Single's Day Global Shopping Festival on 11/11 2019, Alibaba with Apache Flink



Real-time
Events



ingest



Continuous
analytics

Stream Processing
Engine (Flink)

append



Live Report /
Dashboard




- 2 billion events/sec
- 3 TB/sec

- Online training
- Real Time Fraud Detection
- Dynamic Pricing
- ...


(sub)second latency

流计算的广泛应用

Application



Network



System



Smart Manufacturing



Auto-Vehicles



Digital Banks

large volume of data streams



Flink



STORM




...

Outline

- Motivation
 - Why stream processing?
 - Why stream processing on modern HW?
- Past Work
- Future Work

The End of Moore's Law



Source: BACKUS: Comprehensive High-Performance Research Software Engineering Approach for Simulations in Supercomputing Systems, 2019

在新硬件上部署流计算

Network

large volume of data streams




my research focus

System

Hardware



...



...

研发目标

Trend – –

- (i) Stream processing is becoming more critical for big data applications;
- (ii) Modern architectures are developing rapidly (CPUs, GPUs)


Vision – –

We need a high-performance stream processing engine (SPE) on modern architectures.


Outline

- Motivation
 - Why stream processing?
 - Why stream processing on modern HW?
- Past Work
 - What is the current limitation?
- Future Work

观察现象



Apache Storm



Apache Flink

6/7 apps perform worse with more CPU cores on a 32-core machine

*Source: Shuhao Zhang, et al., "Revisiting the Design of Data Stream Processing Systems on Multi-Core Processors," 2017 ICDE

深入理解

■ **Computation**

Useful time

▨ **Front-end stalls**


Instruction fetch

▤ **Back-end stalls**


Data fetch

□ **Bad speculation**

Branch
misprediction



(a) Storm



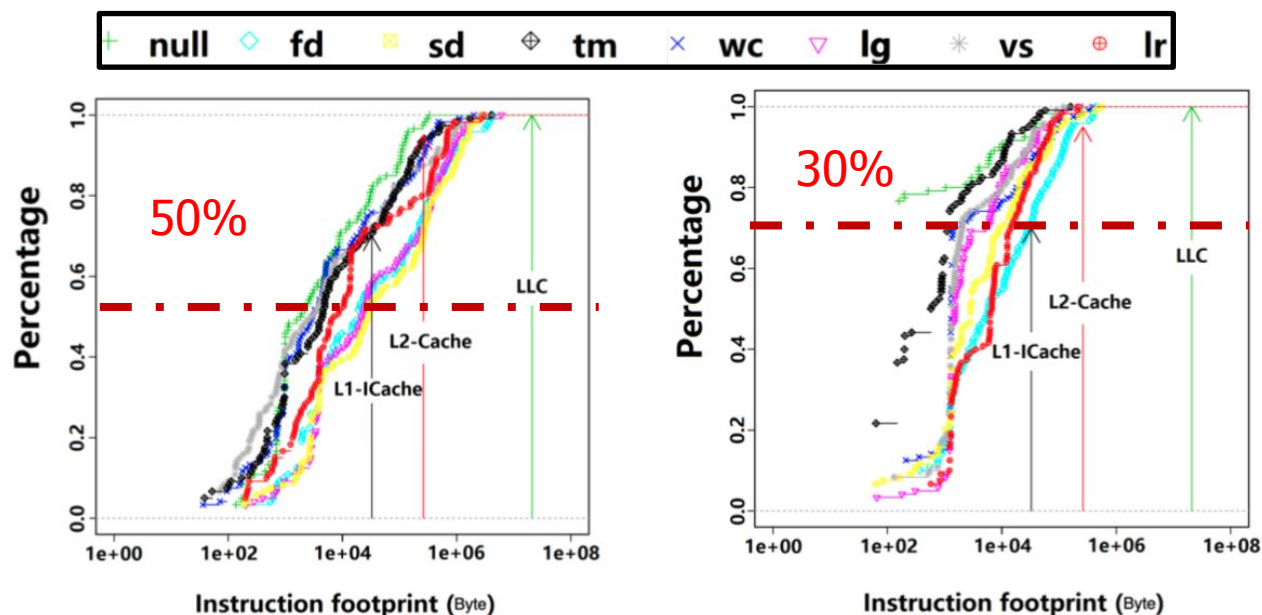
} Processor stalls

(b) Flink

70% of the execution times are wasted

解释现象

L1-ICache size: 32 KB; L2-Cache size: 256KB; L3 Cache size: 20MB



(a) Storm

(b) Flink

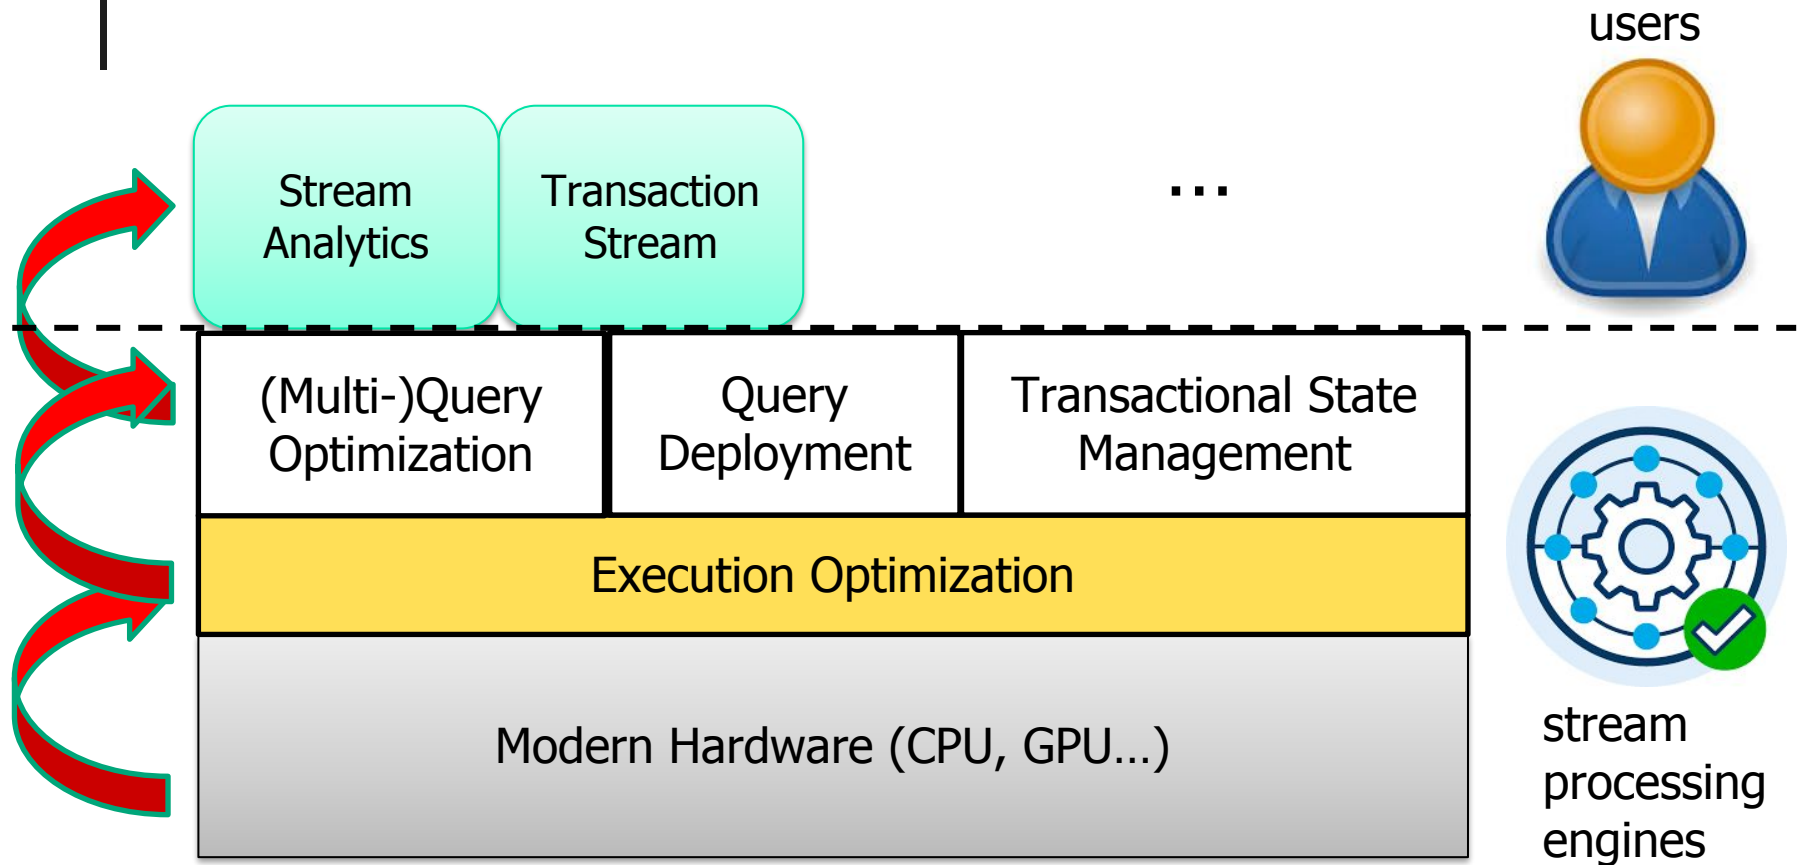
50% instruction footprints \geq instruction cache

Cache trashing

Poor performance

Point (x,y): y% of instruction footprints \leq x

在新硬件上实现高性能流计算




See our ICDE 17 x 2, SIGMOD 19, ICDE 20, ATC 20...

Outline

- Motivation
 - Why stream processing?
- Past Work
 - What is the current limitation?
- **Future Work**

Stream Processing Market Is Fast Growing



Source: "Event Stream Processing Market Research Report
– Global Forecast till 2023" (May. 2019)

The Trend of Stream Processing (1):应用复杂化

- Simple Map-Reduce style stream processing is getting outdated.
- Complex stream processing is getting widely deployed.
 - 实时的机器学习 i.e., Online Learning
 - 流数据挖掘 (e.g., Twitter@KDD19)
 - Live Audio Streaming Services (e.g., Prof.He)
 - Transactional Stream Processing (e.g., S-Store sigmod'15, TStream icde'20)
 - Large Scale Stream Graph Processing (e.g., WuKong+S sosp'17)


The Trend of Stream Processing (2): 需求差异化

- Not just performance.
- But also...
 - Batch + Stream Processing (e.g., Neptune socc'19)
 - Privacy & Security (TimeCrypt, NSDI'20)
 - Unified Computing (e.g., DeltaLake vldb'20)
 - Reproducibility/Deterministic
 - Serverless/Stateful Function
 - ...

The Trend of Stream Processing (3): 硬件多元化

- Specialized stream processing engine is getting more and more important.
 - HW accelerator (e.g., GPU, FPGA, APU etc.)
 - IoT环境 (e.g., NebulaStream)
 - 安全计算硬件 (e.g., StreamBox-TZ)
 - 持久内存 (e.g., Stream Processing with state storing in NVRAM)
 - ...

机遇与挑战共存



Publication Summary

- Patens

First authored patents registered in the USA x2

- Journal

IEEE TPDS 16 x2, ACM SIGMOD Rec. 19 ...

- Conference (only top-tier is shown)

VLDB 13, SC 16, ICDE 17 x2, SIGMOD 19, ICDE 20, IJCAI 20, USENIX ATC 20, SIGMOD 21 (in revision) ...

- System demo and Workshop

VLDB 13, IEEE BIGMM 19, VLIoT 20

- Code distribution

- OmniDB
- BriskStream
- TStream



My Profile Page



ISTD TESTIMONIALS

Shuhao Zhang

Assistant Professor

Website: <https://shuhaozhangtony.github.io/>

Pillar / Cluster: Information Systems Technology and Design

Research Areas: Secure Computing, Artificial and Augmented Intelligence, Data Science, Networked and Autonomous Systems



I'm looking for phd students.

BIOGRAPHY

EDUCATION AND WORK EXPERIENCE

PROFESSIONAL ACTIVITIES

SELECTED KEY PUBLICATIONS

Biography

Shuhao Zhang will be joining ISTD Pillar of SUTD as a faculty in early 2021. He is currently a **postdoctoral researcher** at Technische Universität Berlin working with Prof. Volker Markl on the topic of data stream processing and management. He obtained his Ph.D. degree from the school of computing at National University of Singapore in 2019, under the supervision of Prof. Bingsheng He. He had previously worked as Research Scientist/Assistant with Prof. Bingsheng He, Dr. David B. Lasecki and Dr. Yixin Mi at NUS, SAP and IBM. He had published a list of papers in top-tier

<https://istd.sutd.edu.sg/people/faculty/shuhao-zhang>