

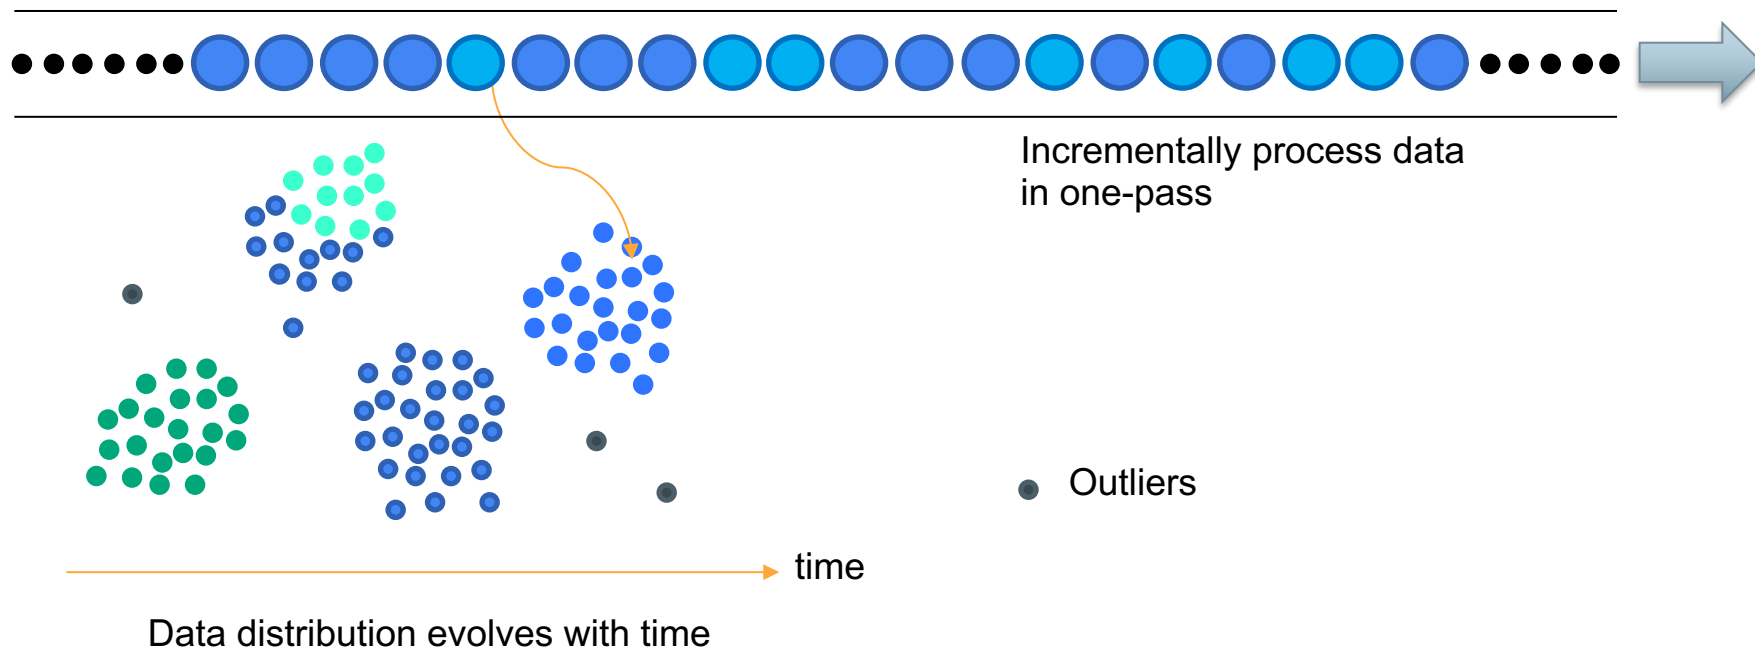
# Data Stream Clustering: An In-depth Empirical Study

Shuhao Zhang (SUTD)

Collaborating with HUST and Sichuan U


# Background: Data Stream Clustering (DSC)

- **Definition:** Partitioning streaming data into clusters in real time.






# Background: Data Stream Clustering (DSC)

- **Definition:** Partitioning streaming data into clusters in real time.
- **Challenges:**
  - Memory limitation ↔ Unbounded data streams
  - Fast response time ↔ Usually expecting fast responses
  - Handle evolving activities
    - Cluster Evolution ↔ Shifting nature of data distributions and the emergence of new outliers over time
    - Outlier Evolution



# Background: Data Stream Clustering (DSC)

- **Definition:** Partitioning streaming data into clusters in real time.
- **Challenges:**
  - Memory Limitation  Unbounded data streams
  - Fast Response Time  Usually expecting fast responses
  - Handle evolving activities
    - Cluster Evolution  Shifting nature of data distributions and the emergence of new outliers over time
    - Outlier Evolution
- **Algorithms:** CluStream [VLDB'03], D-Stream [KDD'07], DBStream [TKDE'16], EDMStream [VLDB'17], SL-Kmeans [NIPS'20] ...

# Background: Literature Gaps and Contributions

- **Coarse-grained comparison:**

- Ignore analysing the impact of individual design aspect of algorithms.



(×) It is not interesting to compare Huawei Mate60 against Iphone14 entirely.

(√) It is more insightful to go into the details: compare their CPU, GPU, memory, etc.

- **Problematic benchmark settings:**

- Not implement algorithms in a unified framework (e.g., varying programming languages, compilers)
- Lack evaluating processing efficiency.



(×) It is less meaningful to compare two algorithms that are implemented in python and C++, respectively.

(√) It is far more fair to implement all algorithms in one language, e.g., C++, before the comparison.



(×) Existing studies only evaluate accuracy.

(√) Both accuracy and efficiency shall be evaluated carefully.

# Background: Our Contributions

## Our Contributions:


- Implement algorithms and designs in an open-sourced platform named **Sesame** (supporting **22** different DSC algorithms, **~13,000** lines of code in C++)
- Evaluate both accuracy and efficiency impact on every single design aspect (our study is like using a “microscope” to observe and analysis DSC algorithms)
- Propose a new algorithm **Benne** through combining flexible design choices achieving **either SoTA accuracy and efficiency** (this is an unexpected gift).

Published (SIGMOD'23)


So, how does a DSC  
algorithm look like actually?




# Design Aspects:



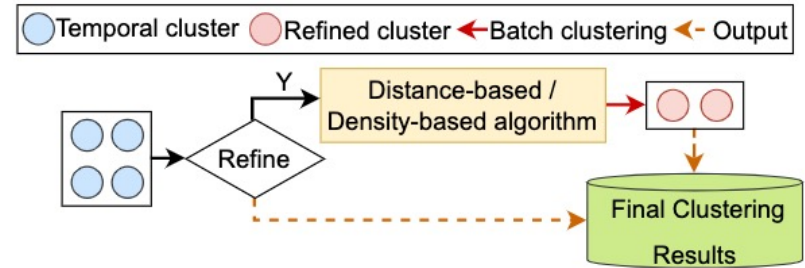
(a) Summarizing Data Structure



(b) Window Model



(c) Outlier Detection




(d) Refinement Strategy


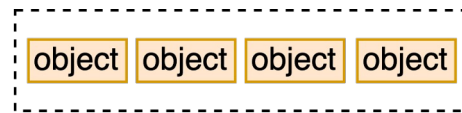


# Design Aspects: Summarizing Data Structure

- **Hierarchical:** Organize temporal clusters into a tree.
  - **Clustering Feature Tree (CFT):** flexibly make adjustment.
  - **Coreset Tree (CoreT):** lazily rebuild the whole structure.
  - **Dependency Tree (DPT):** cluster based on the evolving cluster density.
- **Partitional:** Organize temporal clusters into a list.
  - **Micro Clusters (MCs):** similar structure as *CFT* but under different catalog.
  - **Grids (Grids):** free from frequent distance calculation.
  - **Augmented Meyerson Sketch (AMS):** Frequently reconstruct temporal clusters to keep its total number fixed.



(a) Hierarchical



(b) Partitional

Figure: Two Catalogs of Summarizing Data Structures.

# Design Aspects: Window Model

- **Landmark (*LandWM*)**: Cluster data between two landmarks into a window.
- **Sliding (*SlidingWM*)**: Cluster data whose timestamp falls within current window range.
- **Damped (*DampedWM*)**: Associate data with weights decaying over time.




Figure: Three Types of Window Models.

# Design Aspects: Outlier Detection Mechanism

- **No Outlier Detection (*NoOutlierD*)**
- **Outlier Detection (*OutlierD*):**  
periodically discard sparse clusters
- **Outlier Detection with buffer**  
(*OutlierD-B*): store the sparse clusters into a buffer rather than discarding.
- **Outlier Detection with Timer**  
(*Outlier-T*): additionally check the activity of the discarded clusters
- **Outlier Detection with Buffer and Timer (*OutlierD-BT*)**

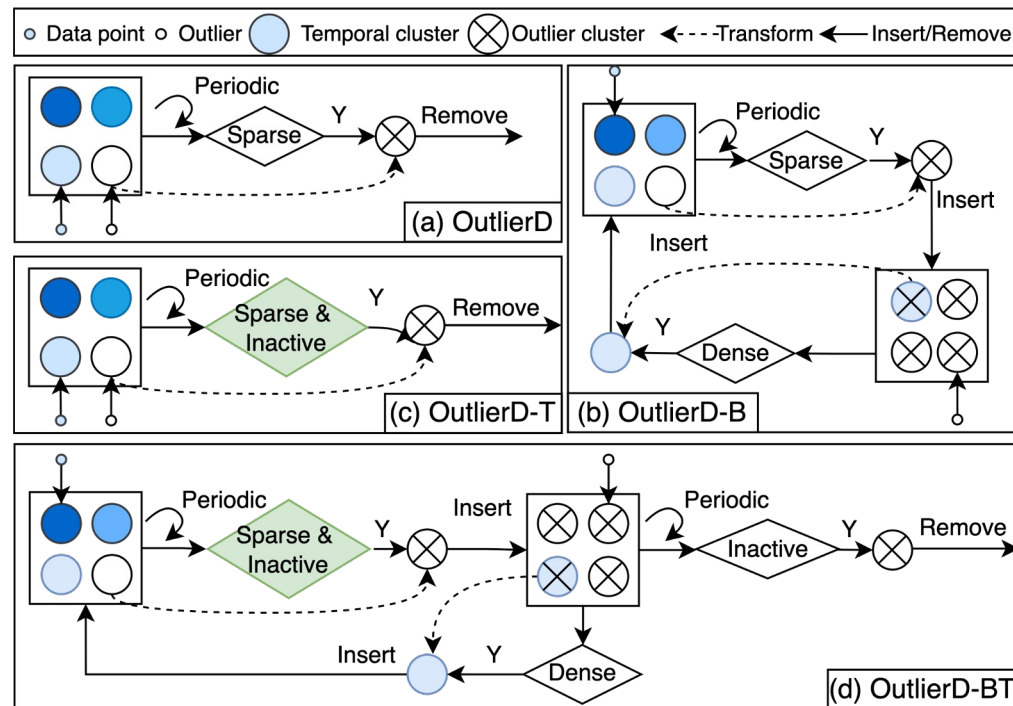


Figure: Four Types of Outlier Detection Mechanism

# Design Aspects: Refinement Strategy

- **With refinement (Refine):** Apply batch clustering algorithms such as KMeans or DBSCAN to further refine the online results before output.
- **Without refinement (NoRefine):** directly output the online temporal clusters as the final clustering results.

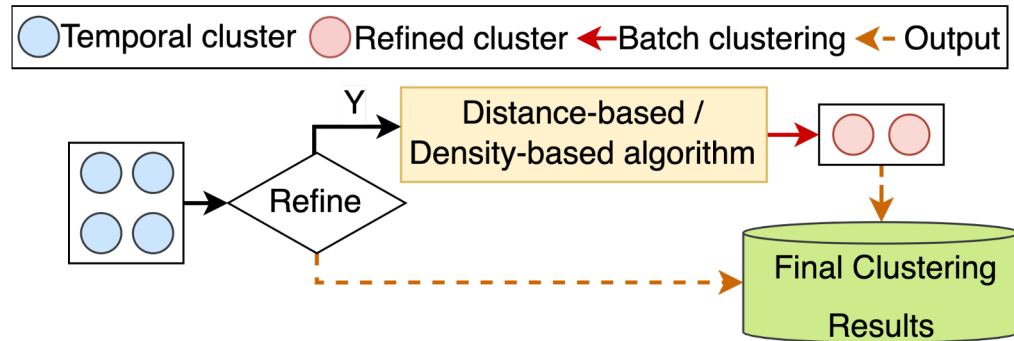


Figure: General workflow of refinement strategy

How do different design  
options really matter?



# Methodology

- Benchmark Testbed: Sesame

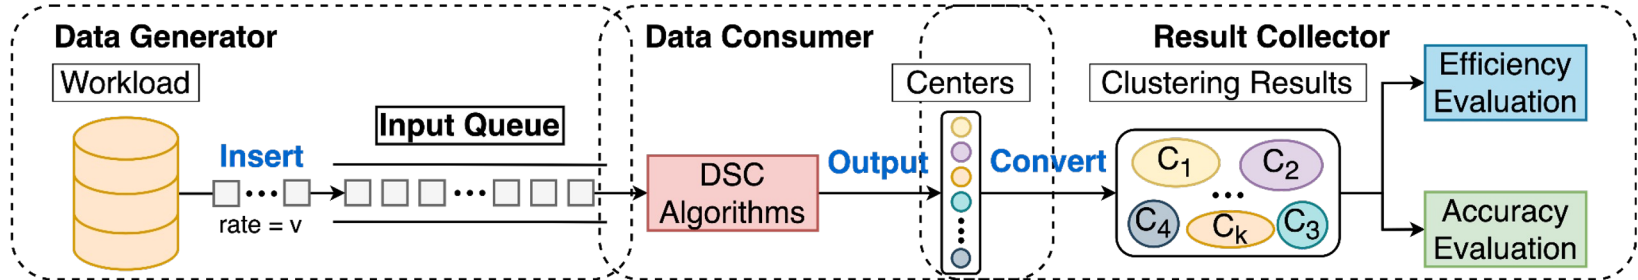


Figure: Sesame Workflow

# Methodology

- Algorithm Selection:

- Cover a wide range of design decisions of all four design aspects.
- Either representative or recently covering a long history in this field.

Algorithm	Year	Summarizing Data Structure		Window Model	Outlier Detection	Offline Refinement
		Name	Catalog			
<u>BIRCH</u>	1996	<i>CFT</i>	<i>Hierarchical</i>	<i>LandmarkWM</i>	<i>OutlierD</i>	<i>NoRefine</i>
<u>CluStream</u>	2003	<i>MCs</i>	<i>Partitional</i>	<i>LandmarkWM</i>	<i>OutlierD-T</i>	<i>Refine</i>
<u>DenStream</u>	2006	<i>MCs</i>	<i>Partitional</i>	<i>DampedWM</i>	<i>OutlierD-BT</i>	<i>Refine</i>
<u>DStream</u>	2007	<i>Grids</i>	<i>Partitional</i>	<i>DampedWM</i>	<i>OutlierD-T</i>	<i>Refine</i>
<u>StreamKM++</u>	2012	<i>CoreT</i>	<i>Hierarchical</i>	<i>LandmarkWM</i>	<i>NoOutlierD</i>	<i>Refine</i>
<u>DBStream</u>	2016	<i>MCs</i>	<i>Partitional</i>	<i>DampedWM</i>	<i>OutlierD-T</i>	<i>Refine</i>
<u>EDMStream</u>	2017	<i>DPT</i>	<i>Hierarchical</i>	<i>DampedWM</i>	<i>OutlierD-BT</i>	<i>NoRefine</i>
<u>SL-KMeans</u>	2020	<i>AMS</i>	<i>Partitional</i>	<i>SlidingWM</i>	<i>NoOutlierD</i>	<i>NoRefine</i>

Table: Selected Algorithm Summary

- **Workload Selection:** [Table: Workload Summary](#)

Workload	Length	Dimension	Cluster Number	Outliers	Evolving Frequency
FCT	581012	54	7	False	Low
KDD99	4898431	41	23	True	Low
Insects	905145	33	24	False	Low
Sensor	2219803	5	55	False	High
EDS	245270	2	363	False	Varying
ODS	100000	2	90	Varying	High


- **Evaluation Metrics:**

- **Accuracy:** We use purity to measure the general clustering quality and also use CMM to test the design aspects' ability to handle cluster evolution.
- **Efficiency:** We use throughput for efficiency comparison.



# Experimental Analysis: Findings and Observations

- **Key Finding 1:** For each design aspect, none of the design choices can always guarantee good performance under varying workload characteristics and/or optimization targets.




**Figure:** Comparison of the Ability for Summarizing Data Structures to Handle Cluster Evolution.

**Observation 1:** *MCs* and *CFT* guarantee high accuracy while *CFT* and *Grids* guarantee high efficiency for handling cluster evolution than other types of data structure.

# Experimental Analysis: Findings and Observations

- **Key Finding 1:** For each design aspect, none of the design choices can always guarantee good performance under varying workload characteristics and/or optimization targets.




**Figure:** Comparison of the Ability for Window Models to Handle Outlier / Cluster Evolution.

**Observation 2:** The efficiency of *LandmarkWM* and *DampedWM* becomes worse with the increase of cluster evolution frequency.

# Experimental Analysis: Findings and Observations

- **Key Finding 1:** For each design aspect, none of the design choices can always guarantee good performance under varying workload characteristics and/or optimization targets.

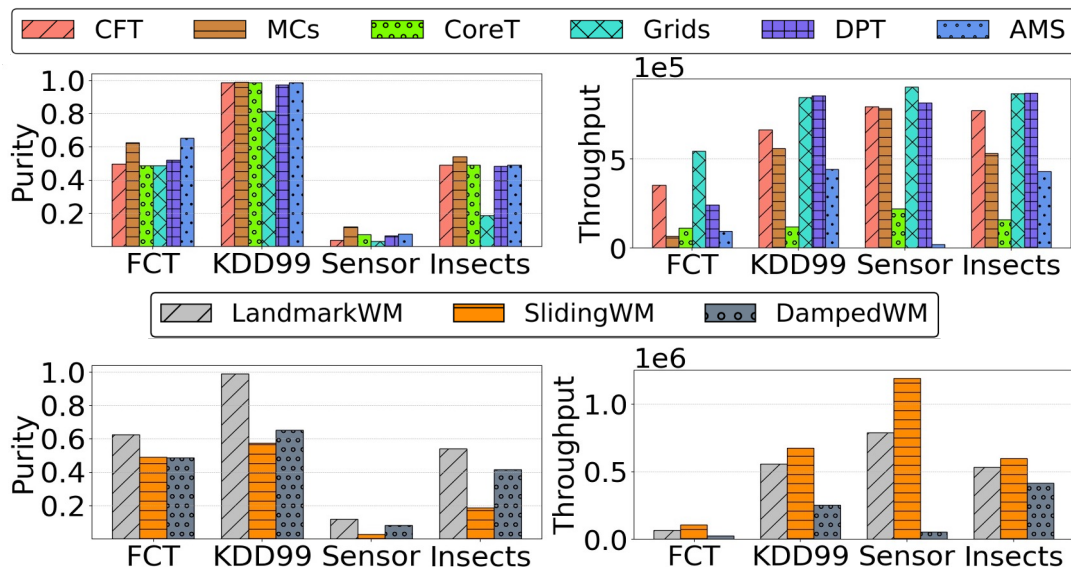


**Figure:** Comparison of the Ability for Outlier Detection Mechanisms to Handle Outlier Evolution.

**Observation 3:** Utilizing a *timer* in outlier detection can greatly improve the clustering accuracy and even increase the efficiency under outlier evolution.

# Experimental Analysis: Findings and Observations

- **Key Finding 2:** Each combined selection of design choices from four design aspects has its own strength and limitation and none can achieve the highest accuracy and efficiency at the same time.



## Observation 3:

{*Grids* summarizing data structure}  
+ {*SlidingWM* window mode}  
lead to high clustering efficiency but low accuracy.

Figure: General Comparison of Summarizing Data Structures / Window Models

# Experimental Analysis: Findings and Observations

- **Key Finding 2:** Each combined selection of design choices from four design aspects has its own strength and limitation and none can achieve the highest accuracy and efficiency at the same time.

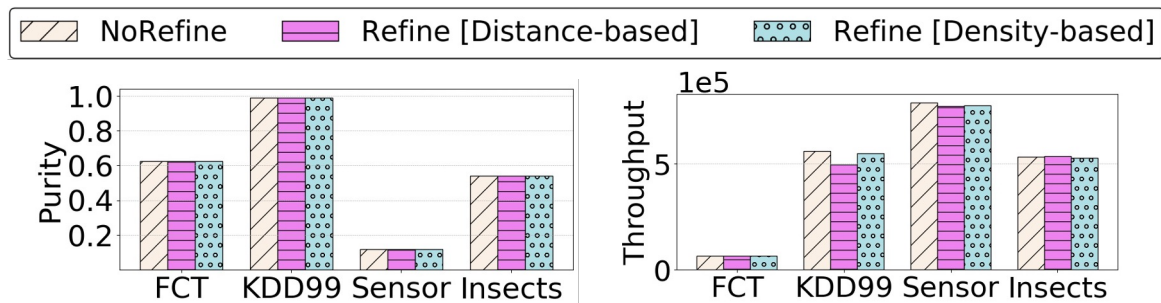


Figure: General Comparison of Offline Refinement Strategy.

**Observation 4:** Applying an *offline refinement strategy* has little impact on both clustering accuracy and efficiency.

**Observation 5:** Composing suitable design choices from each design aspect, we obtain a novel DSC algorithm (i.e., *Benne*) that can be reconfigured to achieve either the highest accuracy or highest efficiency, but not at the same time. (Show later)

# Experimental Analysis: Findings and Observations

- **Key Finding 3:** Algorithm configuration and correlations among design aspects bring further complex influence on the clustering behaviour.

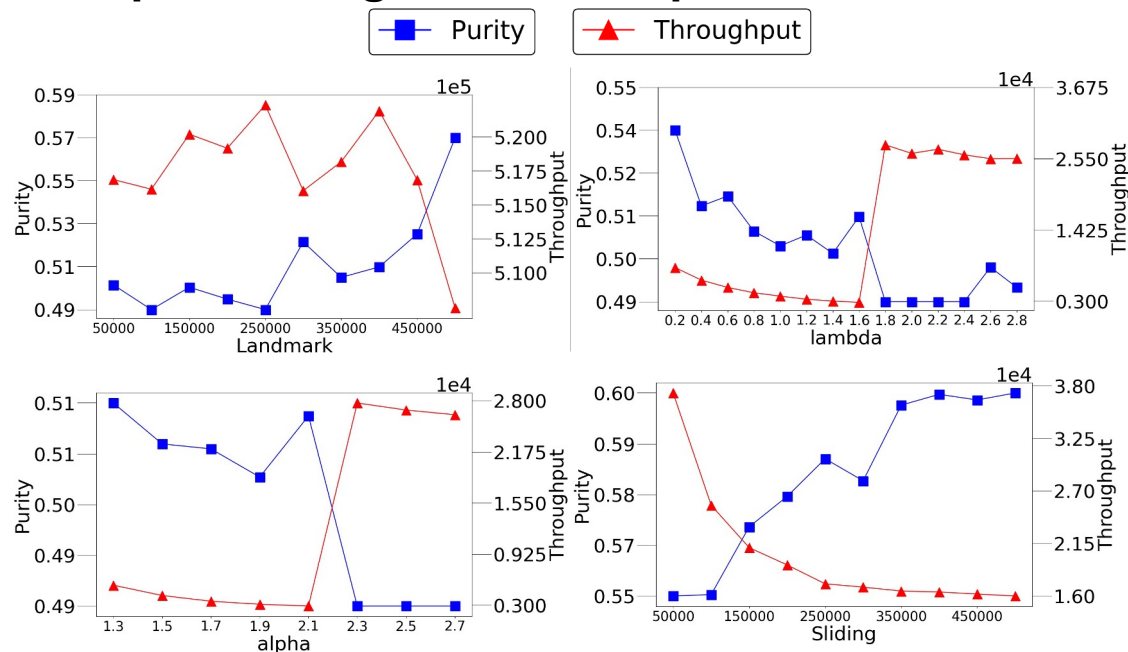


Figure: Configuration Analysis of Window Models.

**Observation 6:** Different algorithm configurations bring *non-trivial trade-offs* in terms of accuracy and efficiency.

**Observation 7:** There is an unsuitable summarizing data structure *dominantly leads to poor performance* of DSC algorithms regardless of the selection of other design choices.

How does the empirical  
analysis useful?



# Experimental Analysis: *Benne* Algorithm

## Algorithm 1: Execution flow of *Benne*.

```
Data:  $p$  // Input point
Data:  $s$  // Summarizing data structure
Input:  $struc.$  // Selected type of summarizing data structure
Input:  $win.$  // Selected type of window model
Input:  $out.$  // Selected type of outlier detection mechanism
Input:  $ref.$  // Selected type of refinement strategy

// Online Phase
1 while !stop processing of input streams do
2   Window Fun. (...);
3   if  $out. \neq \text{NoOutlierD}$  then
4      $b \leftarrow$  Outlier Fun. (...);
5     if  $b = \text{false}$  then
6       Insert Fun. (...) // Insert  $p$  to  $s$  and update  $s$ 
7   else
8     Insert Fun. (...) // Insert  $p$  to  $s$  and update  $s$ 

// Offline Phase
9 if  $ref. \neq \text{NoRefine}$  then
10  Refine Fun. ( $ref.$ );
```

- **Benne (Accuracy):**

***MCs + LandmarkWM + OutlierD-B + NoRefine***

- **Benne (Efficiency):**

***CFT + LandmarkWM + OutlierD-T + NoRefine***



# Experimental Analysis: Overall Comparison

- **Benne (Accuracy)** achieves the best purity.
- **Benne (Efficiency)** achieves the highest throughput.




Figure: General Comparison of Existing DSC algorithms and *Benne*.

- Code, data and scripts are available at <https://github.com/intellistream/Sesame>
- It is purely written in modern C++.
- One can easily reproduce all of our experimental results by “one-click” of our scripts in your machine.
- For easier access by the ML community, we have additionally built a Python API to our framework.

```
1  import benne
2
3  # Create an instance of Parameters
4  params = benne.Parameters()
5
6  # Get and set the algorithm
7  algorithm = params.algo
8
9  print("Current algorithm:", algorithm)
10 params.algo = benne.AlgoType.BIRCH
11 params.input_file = "/home/shaun/Sesame/benchmark/datasets/CoverType.txt"
12 print("Updated algorithm:", params.algo)
13
14 # ... Continue getting and setting other parameters
15 # Accessing docstring
16 print(benne.Parameters.algo.__doc__)
17
18 print("Input file:", params.input_file)
19
20 # Run the SESAME algorithm
21 benne.run()
```

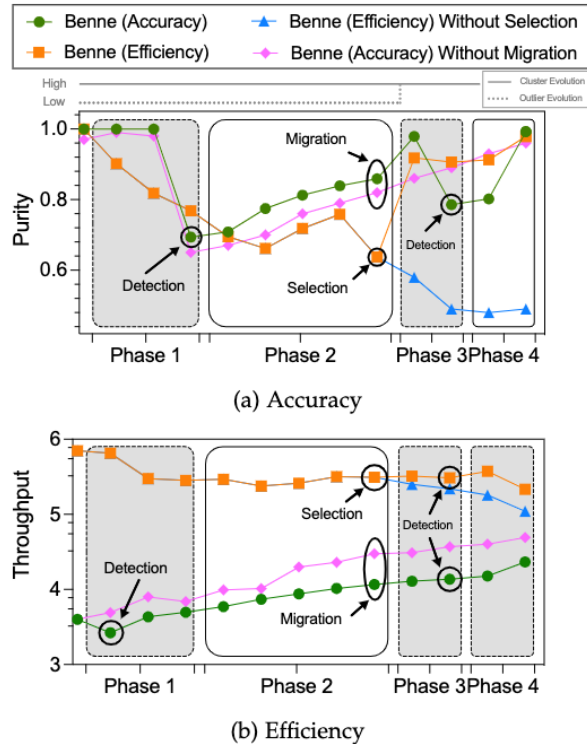
And then?



- **Key Enhancement 1:** Regular Stream Characteristics Detection
- **Key Enhancement 2:** Automatic Design Choice Selection
- **Key Enhancement 3:** Flexible Algorithm Migration

Under Review (TKDE)

# Enhancements



- 1) Both Benne (Accuracy) and Benne (Efficiency) swiftly recover from workload changes.
- 2) Automatic design choice selection is a critical component to ensure the adaptivity.
- 3) Algorithm migration improves accuracy at the expense of clustering speed.

**Figure:** Detailed performance analysis on KDD99 workload

# Enhancements

- 1) For Benne (Accuracy), the time allocation for both detection and migration is relatively minimal in comparison to the primary clustering task.
- 2) For Benne (Efficiency), the proportion of time spent on the detection appears to be larger. However, it's crucial to note that Benne (Efficiency) omits the migration procedure altogether.

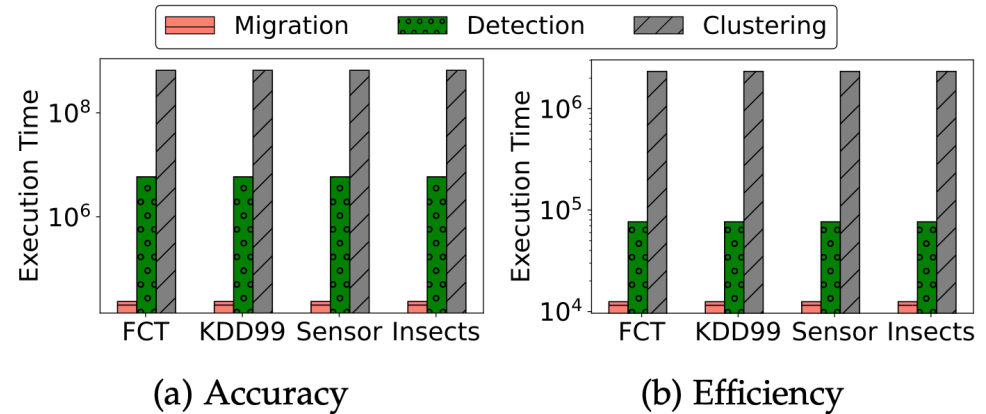


Figure: Execution Time Break Down Analysis

# Future Work



- High performance scalable data stream clustering algorithms (e.g., GPU acceleration)
- High-dimensional data stream clustering (e.g., VectorDB, trajectory data stream analytics)
- Online continual learning (e.g., fast coresets selection)



# Conclusion



# Conclusion

- It is still challenging to balance the trade-off between accuracy and efficiency of DSC algorithms.
- Each design choice has its own pros and cons and should be dynamically adjusted and carefully combined to obtain the best DSC algorithm under different workload characteristics.
- A dynamic algorithm configuration strategy is required for the stream setting.
- Thank you for listening
- Email to ask follow-up questions: [shuhao\\_zhang@sutd.edu.sg](mailto:shuhao_zhang@sutd.edu.sg)